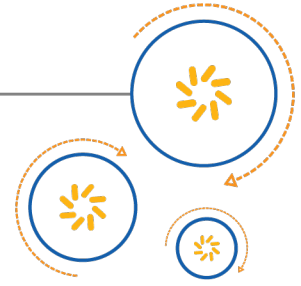




Qualcomm Technologies International, Ltd.



BlueCore Kalimba Architecture 5

User Guide

80-CT525-1 Rev. AC

October 24, 2017

Confidential and Proprietary – Qualcomm Technologies International, Ltd.

NO PUBLIC DISCLOSURE PERMITTED: Please report postings of this document on public servers or websites to DocCtrlAgent@qualcomm.com.

Restricted Distribution: Not to be distributed to anyone who is not an employee of either Qualcomm Technologies International, Ltd. or its affiliated companies without the express approval of Qualcomm Configuration Management.

Not to be used, copied, reproduced, or modified in whole or in part, nor its contents revealed in any manner to others without the express written permission of Qualcomm Technologies International, Ltd.

Qualcomm BlueCore, Qualcomm Kalimba, and Qualcomm aptX are products of Qualcomm Technologies International, Ltd. Other Qualcomm products referenced herein are products of Qualcomm Technologies International, Ltd.

Qualcomm is a trademark of Qualcomm Incorporated, registered in the United States and other countries. BlueCore and aptX are trademarks of Qualcomm Technologies International, Ltd., registered in the United States and other countries. Kalimba is a trademark of Qualcomm Technologies International, Ltd. Other product and brand names may be trademarks or registered trademarks of their respective owners.

This technical data may be subject to U.S. and international export, re-export, or transfer ("export") laws. Diversion contrary to U.S. and international law is strictly prohibited.

Qualcomm Technologies International, Ltd. (formerly known as Cambridge Silicon Radio Limited) is a company registered in England and Wales with a registered office at: Churchill House, Cambridge Business Park, Cowley Road, Cambridge, CB4 0WZ, United Kingdom.
Registered Number: 3665875 | VAT number: GB787433096

Document History

Revision	Date	Change Reason
1	19 FEB 15	Original publication of this document. Alternative document number CS-00318059-UG.
2	27 SEP 16	QTIL brand update and editorial changes
AC	OCT 2017	Document Reference Number updated to Agile number. No change to technical content.

Contents

Document History	2
1 Introduction	10
2 Key Features	12
3 System Overview	14
3.1 Kalimba DSP Core	15
3.2 Kalimba DSP Memory	15
3.3 Kalimba DSP Peripherals	16
3.3.1 Memory Management Unit Interface	16
3.3.2 Programmable I/O Control	16
3.3.3 Interrupt Control	16
3.3.4 Clock Source Select and Timers	16
3.3.5 Debug Interface	16
3.4 Differences in Kalimba Architecture 5 and Architecture 3	16
3.4.1 Interrupt Controller	16
3.4.2 Program Memory Addressing	17
4 Kalimba DSP Core Architecture	18
4.1 Arithmetic Logic Unit	18
4.2 Address Generators	19
4.3 Registers	20
4.4 Bank 1 Registers	20
4.5 rFlags Register	21
4.5.1 Negative Flag (N)	22
4.5.2 Zero Flag (Z)	22
4.5.3 Carry Flag (C)	22
4.5.4 Overflow Flag (V)	23
4.5.5 Sticky Overflow Flag (SV)	23
4.5.6 User Definable Flag (UD)	23
4.5.7 Bit Reverse Flag (BR)	23
4.5.8 User Mode Flag (UM)	23

4.5.9 rFlags Behaviour During Interrupts	23
4.5.10 Condition Codes	24
4.6 rMAC and rMACB Registers	25
4.7 Bank 2 Registers	25
4.7.1 Index Registers	26
4.7.2 Modify Registers	26
4.7.3 Length Registers	26
4.8 Bank 3 Registers	26
4.8.1 Base Registers	27
4.9 Stalls	28
4.10 Prefixes	28
4.11 Circular Buffers	29
4.11.1 Circular Buffers Without Use of Base Registers	29
4.11.2 Circular Buffers With Use of Base Registers	30
4.12 Zero Overhead Looping: do loop	30
4.13 Debug	31
5 Memory Organisation	32
5.1 Memory Map	35
5.1.1 PM Memory Map	35
5.1.2 DM1 Memory Map	35
5.1.3 DM2 Memory Map	36
6 Instruction Set	37
6.1 ADD and ADD with CARRY	38
6.2 SUBTRACT and SUBTRACT With Borrow	40
6.3 Bank 1/2 Register Operations: ADD and SUBTRACT	41
6.4 Logical Operations: AND, OR and XOR	42
6.5 Shifter: LSHIFT and ASHIFT	43
6.6 LSHIFT and ASHIFT (56 bit)	45
6.7 rMAC(B) Move Operations	46
6.8 Multiply: Signed 24-bit Fractional and Integer	47
6.9 MULTIPLY and ACCUMULATE (56 bit)	49
6.10 MULTIPLY and ACCUMULATE (56 bit) with ADD/SUB	51
6.11 LOAD/STORE with Memory Offset	53
6.12 Extended LOAD/STORE with Memory Offset	54
6.13 Sign Bits Detect and Block Sign Bits Detect	55
6.14 Divide Instruction	56
6.15 Stack Instructions	58

6.16 Program Flow: CALL, JUMP, RTS, RTI, and DO...LOOP	60
6.17 Type A Miscellaneous One and Two Operand Instructions	61
6.18 Indexed MEM_ACCESS_1 and MEM_ACCESS_2	63
7 Instruction Coding	65
7.1 Type A Instruction	66
7.2 Type B Instruction	66
7.3 Type C Instruction	66
7.4 Special Cases	67
7.5 OP_CODE Coding	67
7.6 AM Field	69
7.7 Carry Field (C Field)	70
7.8 Bank 1/2 Register Select Field (B2RS Field)	70
7.9 Saturation Select Field (V Field)	70
7.10 Sign Select Field (S Field)	71
7.11 k_{16} Coding for LSHIFT and ASHIFT	71
7.12 rMAC Sub-registers	71
7.13 ASHIFT	72
7.14 LSHIFT	72
7.15 Type A and Type C LSHIFT56 and ASHIFT56	73
7.16 k_{16} Coding Divide Instructions	73
7.17 Type A Miscellaneous One and Two Operand Instruction Encodings	74
7.18 Type A Stack Encoding	74
7.19 Type B Stack Encoding	75
7.20 Type C Stack Encoding	76
7.21 Type C MAC with ADD/SUB Instruction	76
7.22 Type A and B MAC and Optional ADD/SUB Instruction	78
8 Kalimba DSP Peripherals	79
8.1 MMU Interface	80
8.1.1 Read Ports	80
8.1.2 Write Ports	80
8.2 DSP Timers	80
8.3 Kalimba Interrupt Controller	81
8.3.1 Interrupt Controller Functionality	81
8.3.2 Registered Interrupt Processing	81
8.3.3 Interrupt Priorities	82
8.3.4 DSP Core Functionality During Interrupt	82
8.4 Generation of MCU Interrupt	82

8.5 PIO Control from the Kalimba DSP	83
8.6 MCU Memory Windows in DM2	83
8.7 Nonvolatile Memory Windows in DM2	83
8.8 PM Window in DM1	83
8.9 PM Nonvolatile Memory Window with Direct-mapped Cache	83
8.10 Clock Rate Divider Control	84
8.11 Debugging	84
9 Document References	85
10 Terms and Definitions	86
A Number Representation	88
A.1 Binary Integer Representation	88
A.2 Binary Fractional Representation	88
A.3 Integer Multiplication (Signed)	89
A.4 Fractional Multiplication (Signed)	90
B DSP Memory Mapped Registers	91
B.1 DSP Memory Mapped I/O	91
C Software Examples	92
C.1 Double-Precision Addition	92
C.2 Fractional Double-Precision Multiply	92
C.3 Integer Double-Precision Multiply	93
C.4 FIR Filter	93
C.5 Cascaded Bi-Quad IIR Filter	94
C.6 Radix-2 FFT	95
D Bit-reversed Addressing Explained	100

Tables

Table 1-1: Kalimba Architecture Development.....	11
Table 4-1: Bank 1 Registers.....	20
Table 4-2: rFlags Register.....	21
Table 4-3: Condition Codes.....	24
Table 4-4: Bank 2 Registers.....	25
Table 4-5: Bank 3 Registers.....	26
Table 4-6: Prefix Requirement Examples.....	29
Table 5-1: PM Memory Map.....	35
Table 5-2: DM1 Memory Map	35
Table 5-3: DM2 Memory Map.....	36
Table 6-1: Instruction Set Descriptions.....	37
Table 7-1: OPCODE Coding Format.....	67
Table 7-2: AM Field.....	69
Table 7-3: Field Options.....	70
Table 7-4: B2RS Field.....	70
Table 7-5: V Field.....	70
Table 7-6: S Field.....	71
Table 7-7: k_{16} Coding Shift Format.....	71
Table 7-8: rMAC Sub-registers.....	71
Table 7-9: ASHIFT.....	72
Table 7-10: LSHIFT.....	72
Table 7-11: Divide Field.....	73
Table 7-12: Divide Field States.....	73
Table 7-13: Encodings for One and Two Operand Instructions	74

Table 7-14: Type A Stack Operation Encoding.....	74
Table 7-15: Type B Stack Operation Encoding.....	75
Table 7-16: BankSelect = 0 (Bank 1), or BankSelect = 1 (Bank 2).....	76
Table 7-17: Type C Stack Operation Encoding.....	76
Table 7-18: BlueCore7/BlueCore8 Type C MAC with ADD/SUB Encoding.....	77
Table 7-19: BlueCore7/BlueCore8 Type A and B MAC and Optional ADD/SUB Encoding.....	78
Table 9-1: Document References.....	85
Table D-1: Binary Representation of 8-Word Buffer.....	100

Figures

Figure 3-1: Kalimba DSP Coprocessor Subsystem..... 15

Figure 4-1: Kalimba DSP Core Base Architecture..... 18

Figure 4-2: rMAC and rMACB Sub-register Definitions..... 25

Figure 5-1: Program Memory Organisation for CSR8675.....33

Figure 5-2: Data Memory Organisation for CSR8675..... 34

Figure 7-1: Instruction Coding Format..... 65

Figure 8-1: Kalimba DSP Peripheral Interfaces..... 79

Figure 8-2: Example of MMU Interface Usage for a Wireless MP3 Player..... 80

Figure A-1: Binary Integer Representation.....88

Figure A-2: Binary Fractional Representation..... 88

Figure A-3: Integer Multiplication..... 89

Figure A-4: Fractional Multiplication (Signed)..... 90

1 Introduction

This user guide is for developers of software applications and algorithms for the fifth-generation Qualcomm® Kalimba™ DSP coprocessor, as implemented on selected Qualcomm Technologies International, Ltd. (QTI) devices.

The guide documents the architecture of the Kalimba DSP, instruction set description and peripheral features, and includes some example code. Read this document in conjunction with the other Kalimba DSP tools documents that are available.

Kalimba DSP is present in many QTI devices. The Kalimba Architecture 5 DSP (KalArch5) particularly targets audio processing applications. Example audio processing applications include:

- SBC encoding and decoding, as defined in the Bluetooth A2DP
- MP3 encoding and decoding, as defined in ISO/IEC 11172-3, and the sample rate extensions defined in ISO/IEC 13818-3
- AAC encoding and decoding, as defined in ISO/IEC 13818-7
- Qualcomm® aptX™ encoding and decoding
- Alternative voice/Hi-Fi codecs
- Echo and noise cancellation
- Audio signal enhancement:
 - Stereo enhancement
 - Equaliser
 - Lost packet concealment
- Text-to-speech
- Voice recognition

The Kalimba DSP also is suitable for non-audio applications such as picture compression, image processing, communications protocols and general-purpose application code.

Table 1-1 Kalimba Architecture Development

Kalimba Architecture	Products
1 = KalArch1	Qualcomm® BlueCore™ 3-Multimedia
2 = KalArch2	BlueCore5-Multimedia
3 = KalArch3	<ul style="list-style-type: none"> ■ CSR8311 ■ CSR860x ■ CSR861x ■ CSR862x ■ CSR863x ■ CSR864x ■ CSR8670 ■ CSR8810 ■ CSR8811 ■ CSR8820
5 = KalArch5	<ul style="list-style-type: none"> ■ CSR8675 ■ CSRS3703 ^a

^a Based on KAS version of KalArch5, for more information see *KAS Kalimba Architecture 5 DSP User Guide*

2 Key Features

Kalimba DSP core is a 24-bit fixed core with a variety of features to support low-power signal processing.

The key features of the Kalimba DSP core include:

- 24-bit fixed point DSP core
- 120 MHz performance, which can be divided down for power saving
- One program memory and two data memory banks, all three of which can be accessed simultaneously in a single cycle
- Flash/ROM support for both data and code, with caches to improve code performance
- Single-cycle 24 x 24-bit multiply with two 56-bit accumulators
- Single-cycle barrel shifter with 56-bit input and 56-bit or 24-bit output
- Thirteen-cycle divide (performed in the background)
- Majority of instructions can be conditional
- Zero overhead ring buffer indexing
- Zero overhead looping and unconditional branching
- Bit reversed addressing capability, and bit reversed data function
- Largely orthogonal instruction set, which is quick to learn and easy to write in the algebraic assembler language
- **Stack instructions** (PUSH, POP, PUSHM, POPM, FP/SP Adjust, FP/SP Relative LOAD/STORE), and other instructions featuring overflow detection
- Low-power internal architecture
- Eight hardware program breakpoints and two data breakpoints (each covering an address range start to end)

The key features of the Kalimba DSP peripherals include:

- Close integration with the on-chip MMU, giving access to features such as DACs and ADCs
- 12 low-overhead read/write ports to transfer streaming data to and from the BlueCore subsystem
- 2 memory-mapped windows into the MCU RAM for data exchange
- 3 windows for access to the flash/ROM data memory
- Memory-mapped interface to the I/O address map

- Multiple interrupt sources including two 32-bit timers
- Read, write and direction control access to external PIO lines

3 System Overview

All Kalimba Architecture 5 devices contain the Kalimba DSP, which implements a standard set of functional elements.

The architecture consists of the following functional elements:

- Kalimba DSP core
- DSP memory, this RAM is used for:
 - Data memory
 - Program memory
- Memory mapped I/O
- MMU interface
- Programmable I/O control
- Interrupt control
- Clock source

- Timers
- Debug interface

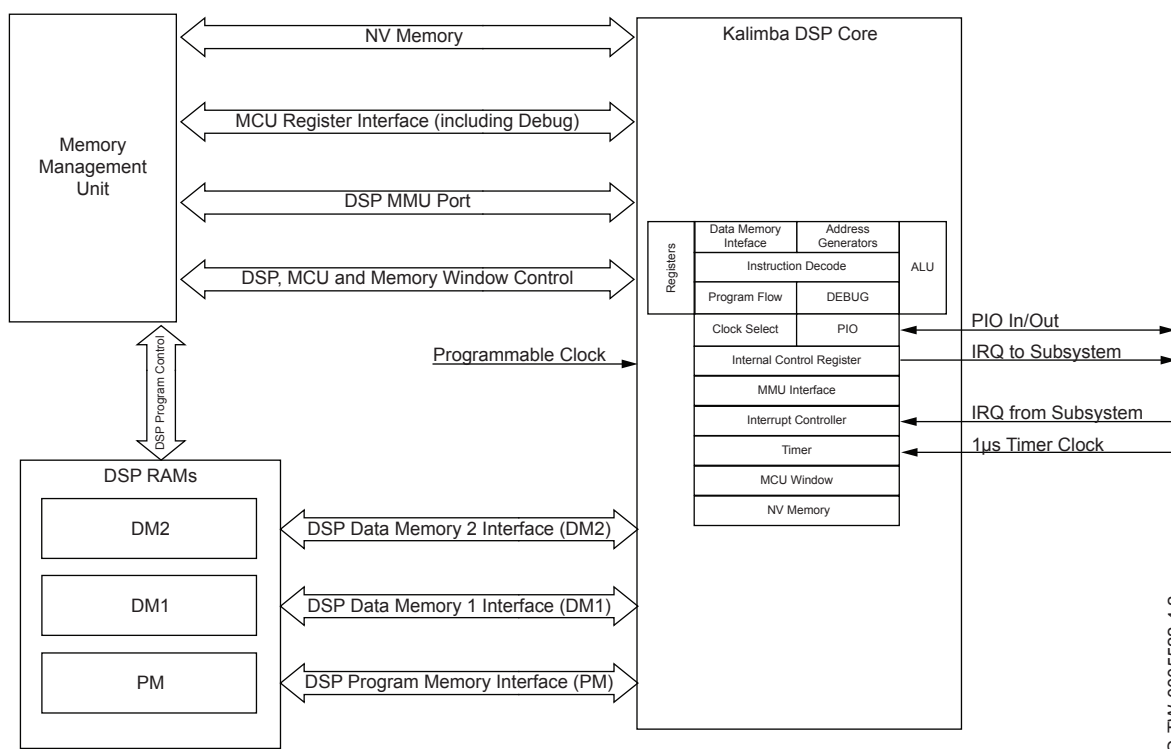


Figure 3-1 Kalimba DSP Coprocessor Subsystem

3.1 Kalimba DSP Core

The Kalimba DSP core is an open-platform DSP that can perform signal processing functions on over-air data or audio codec data to enhance audio applications.

The Kalimba DSP also is available as a general-purpose processor. Figure 3-1 and Section 8 show how the DSP interfaces to other functional blocks within the BlueCore device.

3.2 Kalimba DSP Memory

The Kalimba DSP contains two 24-bit data memories and a 32-bit program memory.

The Kalimba DSP contains the following on-chip RAM:

- DM1: 24-bit data memory 1
- DM2: 24-bit data memory 2
- PM: 32-bit program memory, with 1024-word or 512-word, direct-mapped cache for flash/ROM program space, and 32-word loop cache

Both data and code have a 24-bit address width. Separate documents describe specific product information.

3.3 Kalimba DSP Peripherals

3.3.1 Memory Management Unit Interface

The MMU interface consists of a number of read and write ports that can efficiently transfer streams of data to and from the rest of the IC.

3.3.2 Programmable I/O Control

BlueCore ICs have up to 32 programmable I/O lines controlled by firmware running on the device. The Kalimba DSP core can read any digital I/O directly but can only write to or change the pin direction of digital outputs that the MCU has enabled (carried out through the VM application, if present).

3.3.3 Interrupt Control

The interrupt controller function within the Kalimba DSP provides interrupt control of the Kalimba DSP core. The function allows interrupt sources to select and control three levels of priority settings.

Alongside the interrupts caused by hardware, four software event interrupts are available. Interrupts are serviced between three and five instructions after the interrupt request line goes high (up to and including the automatic branch to the beginning of the interrupt handler).

3.3.4 Clock Source Select and Timers

The Kalimba DSP has a clock source select interface. This clock-rate divider circuit is controllable from the Kalimba DSP core. The Kalimba DSP also has two timers with a 1 μ s time base available.

3.3.5 Debug Interface

The BlueCore device contains a hardware interface that assists in the debugging of applications running on the Kalimba DSP core.

3.4 Differences in Kalimba Architecture 5 and Architecture 3

3.4.1 Interrupt Controller

Kalimba Architecture 5 (KalArch5) interrupt controller has been modified. In earlier devices, each interrupt source could be mapped to three different priority levels, which enabled higher priority interrupt sources to interrupt lower priority sources.

The KalArch5 implementation is slightly different. Each interrupt source is routed into the interrupt controller several times, once for each priority level. For example, the PIO interrupt appears three times.

- \$INT_LOW_PRI_SOURCES_EN_PIO_EVENT_POSN
- \$INT_MED_PRI_SOURCES_EN_PIO_EVENT_POSN
- \$INT_HIGH_PRI_SOURCES_EN_PIO_EVENT_POSN

The functionality is the same as KalArch3, where higher priority interrupts interrupt lower priority ones.

QTIL reference libraries support the new hardware and maintain the current API. Applications that use the libraries are not impacted by these hardware changes.

3.4.2 Program Memory Addressing

KalArch5 addresses program memory in 8-bit words. In previous devices, program memory is addressed in 32-bit words.

This results in the following two changes:

- Program counter increments
- Region sizes in link scripts

The program counter increments differently in KalArch5. In linear code execution, the program counter increments by 4. In earlier devices, the program counter increments by 1.

Regions defined in link scripts are now four times larger. The amount of memory is exactly the same. On the KalArch3, 1 kB of memory is listed as 256 words. On the KalArch5, 1 kB is listed as 1024 words. In both cases, exactly the same amount of memory is used. This means regions defined in the link scripts might appear larger but are the same as in the KalArch3.

As the linker resolves program memory symbols, no code needs updating. After code is reassembled, the linker automatically updates any references. However, if an application uses custom link scripts, these must be updated. The standard link scripts provided by QTIL have been updated.

4 Kalimba DSP Core Architecture

The Kalimba DSP core consists of multiple components, including an ALU, address generators, and a variety of registers.

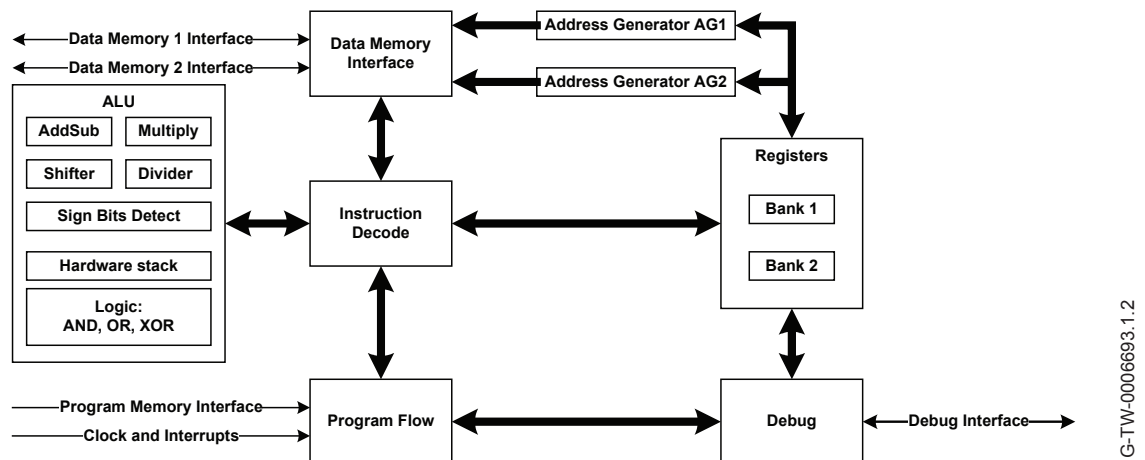


Figure 4-1 Kalimba DSP Core Base Architecture

4.1 Arithmetic Logic Unit

The ALU performs calculations, applies logic, and performs stack instruction control.

The ALU performs the following functions:

- Add and subtract arithmetic
- Logic operations: AND, OR, and XOR
- Single-cycle multiply, multiply/add, and multiply/subtract
- Single-cycle multiply, multiply/add, and multiply/subtract with parallel add and subtract
- Logical and arithmetic shift operations with up to 56-bit output and up to 56-bit input
- Derive exponent and block derive exponent operations, which detect the number of redundant sign bits
- Signed divide, taking a 48-bit dividend (numerator) and a 24-bit divisor (denominator)

NOTE This is a background operation, taking 13 clock cycles.

- Stack instruction control:

□ PUSH

- ☐ POP
- ☐ PLOOK
- ☐ PUSHM
- ☐ POPM

with overflow detection

4.2 Address Generators

Address Generators (AGs) form data memory addresses for indexed memory reads or writes. Each AG has four associated address pointers (index registers).

When an index register is used for a memory access, it is post-modified by a value in a specified modify register, or by a 2-bit constant. With two independent AGs, the DSP can generate two addresses simultaneously for dual-indexed memory accesses.

Length values can be associated with four of the index registers to implement automatic modulo addressing for circular buffers. Start addresses also can be associated, removing the need for circular buffer alignment.

When the appropriate mode bit is set in the rFlags register, the output of AG1 is bit-reversed then driven on to the address bus. This feature enables efficient addressing in Radix-2 FFT algorithms. See the Radix-2 FFT code in [Software Examples](#).

4.3 Registers

There are three banks of 16 registers. Bank 1 is general registers available to virtually all instructions (see [Bank 1 Registers](#)). Bank 2 registers are for the control of index memory accesses (see [rMAC and rMACB Registers](#)). Bank 3 registers can be altered using stack instructions (see [Bank 3 Registers](#)).

4.4 Bank 1 Registers

Bank 1 registers are general registers used by all instructions.

Table 4-1 Bank 1 Registers

No.	Name	No. of bits	Description
0	Null	N/A	Always read as 0, writing only affects flags (so can be used for condition testing)
1	rMAC	24/56	The 56 bits are used for multiply accumulate instructions and the input to shift operations. For 24-bit operations: <ul style="list-style-type: none"> ■ Read as bits [47:24] with saturation and unbiased rounding ^a ■ Written as bits [47:24] with sign extension and trailing 0 padding to make 56 bits
2	r0	24	General register
3	r1	24	General register
4	r2	24	General register
5	r3	24	General register
6	r4	24	General register
7	r5	24	General register
8	r6	24	General register
9	r7	24	General register
10	r8	24	General register
11	r9	24	General register
12	r10	24	General register and is used as the loop counter for zero overhead loops.
13	rLink	24	Call instructions put the return PC address in this register for use by <i>rts</i> instructions. ^b
14	rFlags	24	Status and mode flags. See below for a description.
15	rMACB	24/56	The 56 bits are used for multiply accumulate instructions and the input to shift operations. For 24-bit operations: <ul style="list-style-type: none"> ■ Read as bits [47:24] with saturation and unbiased rounding ^a ■ Written as bits [47:24] with sign extension and trailing 0 padding to make 56 bits

^a Unbiased rounding is:

$rMAC_{rounded} = rMAC[47:24] + rMAC[23];$

if ($rMAC[23:0] == 0x800000$) then $rMAC_{rounded}[0] = 0;$

When `rMAC[23:0]` is exactly at the midpoint (`0x800000`), this has the effect of rounding odd `rMAC[47:24]` values away from zero and even `rMAC[47:24]` values towards zero, yielding a zero large sample bias assuming uniformly distributed values.

- ^b There is no hardware subroutine stack, to enable multi-depth subroutine calls the return PC addresses are usually stored on the general-purpose hardware stack.

NOTE Only registers `rMAC` and `r0` to `r5` can be used by indexed memory access instructions. All registers are set to 0 on a Kalimba reset.

4.5 rFlags Register

The `rFlags` register is a 24-bit register that is located in register bank 1 of Kalimba, see [Table 4-2](#).

[Table 4-2](#) shows the individual bits that make up the `rFlags` register and their value after reset. This register has a natural split into 3 bytes:

- The LS byte contains the active flags used by Kalimba.
- The middle byte contains a stored value of the flags.
- The MS byte has no meaning.

The flags are stored when an interrupt has occurred and then restored after Kalimba has finished servicing the interrupt. The `INT_` versions of the various flags are the copies that are stored at the point of the service interruption. On an interrupt, the LS byte is copied to the middle byte, which stores a copy of the flags. The `rti` instruction then automatically restores the flags to their previous value.

Table 4-2 rFlags Register

Name	Bit	Reset State
<code>INT_UM_FLAG</code>	15	0
<code>INT_BR_FLAG</code>	14	0
<code>INT_SV_FLAG</code>	13	0
<code>INT_UD_FLAG</code>	12	0
<code>INT_V_FLAG</code>	11	0
<code>INT_C_FLAG</code>	10	0
<code>INT_Z_FLAG</code>	9	0
<code>INT_N_FLAG</code>	8	0
<code>UM_FLAG</code>	7	0
<code>BR_FLAG</code>	6	0
<code>SV_FLAG</code>	5	0
<code>UD_FLAG</code>	4	0
<code>V_FLAG</code>	3	0
<code>C_FLAG</code>	2	0
<code>Z_FLAG</code>	1	0
<code>N_FLAG</code>	0	0

4.5.1 Negative Flag (N)

This flag is set if the result of the instruction is negative (that is, if the most significant bit is set). Otherwise, this flag is cleared.

4.5.2 Zero Flag (Z)

This flag is set if the result of the instruction is 0. Otherwise, this flag is cleared.

4.5.3 Carry Flag (C)

The carry flag is set or cleared according to the operation performed and the result of the operation.

The state of the carry flag is determined as follows:

- For an addition, C is set if the addition produced a carry (that is, an unsigned overflow). Otherwise, the flag is cleared.
- For a subtraction, C is cleared if the subtraction produces a borrow (that is, an unsigned underflow). Otherwise, the flag is set.
- For other operations (including multiply accumulate), C is cleared.

4.5.4 Overflow Flag (V)

The overflow flag is set or cleared according to mathematical operations and overflows.

The state of the overflow flag is determined as follows:

- For addition, subtraction, arithmetic shifts, integer multiplies, and multiply accumulates, the flag is set if a signed overflow occurred. Operands and the result are regarded as 2's complement signed integers. Otherwise, the flag is cleared.
- The setting and clearing of the flag for the `rMAC` or `rMACB` registers occurs if there is overflow past the 56th bit. For 24-bit registers, the flag is set or cleared if an overflow occurs past the 24th bit.
- The flag is cleared for other operations.

4.5.5 Sticky Overflow Flag (SV)

This flag is set whenever V is set but SV can be cleared only by software explicitly writing to the `rFlags` register. If V is set, SV also is set if the next instruction is an `nop` or if a `prefix` instruction is executed. For example, SV is not set on a type A instruction where the condition is not satisfied.

4.5.6 User Definable Flag (UD)

A special `USERDEF` condition code is `TRUE` if this flag is set and `FALSE` if this flag is clear. Use this flag in code sections to improve speed and code clarity where a particular instruction must be executed conditionally.

4.5.7 Bit Reverse Flag (BR)

If set, the output of `AG1` (index registers `I0` to `I3`) is bit-reversed before being driven to the address bus. Here, bitreversing applies to the least significant 23 bits. The most significant bit selects the data memory used so it stays in place.

4.5.8 User Mode Flag (UM)

If set, interrupts are serviced. On entry to the interrupt service routine (default `PC` address `0x0002`), this flag is cleared. When cleared, no further interrupts are serviced unless the flag is set manually (for example, to support interrupt priority). Execution of an `rti` instruction sets this flag to the value of `INT_UM_FLAG` (normally set unless altered in software).

4.5.9 rFlags Behaviour During Interrupts

Before entering the interrupts, each element of `rFlags` is copied into its interrupt duplicate. For example, `UM_FLAG` is copied into `INT_UM_FLAG`. With the exception of the `UM_FLAG`, all bits remain unchanged.

`UM_FLAG` is cleared and remains cleared unless software enables it. When returning from an interrupt, using the `rti` instruction, the interrupt bits are copied back, overwriting the noninterrupt values.

4.5.10 Condition Codes

The state of the flags present in the `rFlags` register forms the basis of the condition codes supported by conditional instructions in Kalimba.

Table 4-3 Condition Codes

Condition	Condition Flag State	Condition Code
Z (zero)/EQ (equal)	$Z = 1$	0 0 0 0
NZ (not 0)/NE (not equal)	$Z = 0$	0 0 0 1
C (ALU carry)/NB (not ALU borrow)	$C = 1$	0 0 1 0
NC (not ALU carry)/B (ALU borrow)	$C = 0$	0 0 1 1
NEG (negative)	$N = 1$	0 1 0 0
POS (positive)	$N = 0$	0 1 0 1
V (ALU overflow)	$V = 1$	0 1 1 0
NV (not ALU overflow)	$V = 0$	0 1 1 1
HI (unsigned higher)	$C = 1 \text{ AND } Z = 0$	1 0 0 0
LS (unsigned lower or same)	$C = 0 \text{ OR } Z = 1$	1 0 0 1
GE (signed greater than or equal)	$N = V$	1 0 1 0
LT (signed less than)	$N \neq V$	1 0 1 1
GT (signed greater than)	$Z = 0 \text{ AND } N = V$	1 1 0 0
LE (signed less than or equal)	$Z = 1 \text{ OR } N \neq V$	1 1 0 1
USERDEF (user defined)	$USERDEF = 1$	1 1 1 0
Always true	Don't care	1 1 1 1

4.6 rMAC and rMACB Registers

The `rMAC` and `rMACB` registers are 56-bit registers located in register bank 1 of Kalimba (see [rFlags Register](#))

The `rMAC (B)` registers split into a set of separately accessible sub-registers.

- `rMAC (B)` is the overall 56-bit register.
- `rMAC (B) 0` is a 24-bit register that forms the lower part of the `rMAC (B)` register.
- `rMAC (B) 1` is a 24-bit register that forms the middle part of the `rMAC (B)` register.
- `rMAC (B) 2` is an 8-bit register that forms the higher part of the `rMAC (B)` register.
- `rMAC (B) 12` is a 32-bit register that is a combination of `rMAC (B) 2` and `rMAC (B) 1` that forms part of the `rMAC (B)` register.
- The 24-bit rounded and saturated version of `rMAC (B)` is often referred to as `rMAC (B) 24`.

Figure 4-2 shows the size of the rMACA and rMACB registers.




Figure 4-2 rMAC and rMACB Sub-register Definitions

4.7 Bank 2 Registers

Bank 2 registers control index memory accesses and set up modulo addressing through a circular memory buffer.

Table 4-4 Bank 2 Registers

No.	Name	No. of bits	Description
0	I0	24	Index register for AG1
1	I1	24	Index register for AG1
2	I2	24	Index register for AG1
3	I3	24	Index register for AG1
4	I4	24	Index register for AG2
5	I5	24	Index register for AG2
6	I6	24	Index register for AG2
7	I7	24	Index register for AG2
8	M0	24	Modify register for any index register
9	M1	24	Modify register for any index register
10	M2	24	Modify register for any index register
11	M3	24	Modify register for any index register

Table 4-4 Bank 2 Registers (cont.)

No.	Name	No. of bits	Description
12	L0	24	Length register for Index register I0
13	L1	24	Length register for Index register I1
14	L4	24	Length register for Index register I4
15	L5	24	Length register for Index register I5

NOTE Even though the modify and length registers are 24-bit, only the LS 16 bits are used for indexed memory accesses. That is, circular buffers greater than a 64K word in size are not supported.

All registers are set to 0 on DSP reset.

4.7.1 Index Registers

An index register contains an address pointer to data memory and is for indexed addressing.

These registers enable transfer of data to and from selected Bank 1 registers using the address contained within this register (see Section 6.18). Index registers I0 to I3 are associated with AG1 (address generator 1). Index registers I4 to I7 are associated with AG2 (address generator 2).

4.7.2 Modify Registers

When an index register is used for a memory access, it can be post-modified by a value contained in a modify register, or by a 2-bit constant.

4.7.3 Length Registers

Length register values are associated with the four Bank 2 index registers, L0, L1, L4 and L5 (see Table 4-4). The values are for implementing automatic modulo addressing for circular buffers. To disable automatic modulo addressing, set the corresponding length register to 0.

4.8 Bank 3 Registers

The rMAC(B) subregisters are available in LSHIFT, ASHIFT, and rMAC(B) move instructions (see Section 6.5 and Section 6.7). The DivResult and DivRemainder registers can be accessed using Divide instructions (see Section 6.14). The stack instructions can use the SP and FP registers, which are particularly useful for subroutine entry and exit.

Table 4-5 Bank 3 Registers

No.	Name	No. of bits	Description
0	rMAC2	8	The MS 8 bits of rMAC
1	rMAC12	24	The middle 24 bits of rMAC (with sign extension into rMAC2 when popping) ^a

Table 4-5 Bank 3 Registers (cont.)

No.	Name	No. of bits	Description
2	rMAC0	24	The LS 24 bits of rMAC
3	DoLoopStart	24	Start address for do...loops Identical to memory-mapped MM_DOLOOP_START
4	DoLoopEnd	24	End address for do...loops Identical to memory-mapped MM_DOLOOP_END
5	DivResult	24	Divider result Identical to memory-mapped MM_QUOTIENT
6	DivRemainder	24	Divider remainder Identical to memory-mapped MM_REM
7	rMACB2	8	The MS 8 bits of rMACB
8	rMACB12	24	The middle 24 bits of rMACB (with sign extension into rMACB2 when popping) ^b
9	rMACB0	24	The LS 24 bits of rMACB
10	B0	24	Base register 0 (when non-zero, sets the base address of circular buffers and bit-reverse arrays)
11	B1	24	Base register 1
12	B4	24	Base register 4
13	B5	24	Base register 5
14	FP	24	Frame pointer Identical to memory-mapped FRAME_POINTER
15	SP	24	Stack pointer Identical to memory-mapped STACK_POINTER

^a The source register is in rMAC1. The destination register is in rMAC12. See individual instructions for an explanation of how the top 8 bits of rMAC12 are filled.

^b The source register is in rMACB1. The destination register is in rMACB12. See individual instructions for an explanation of how the top 8 bits of rMACB12 are filled.

NOTE Pushing DivResult or DivRemainder stalls the DSP until any background divide has completed. Popping cancels any current background divide.

All registers are set to 0 on DSP reset.

4.8.1 Base Registers

Base registers, B0, B1, B4, and B5 enable circular buffers to be located at an arbitrary address in memory (see Section 4.11).

Base registers enable the option of easier memory management by removing the constraint on circular buffer placement.

4.9 Stalls

The Kalimba DSP has pre-fetch logic that looks ahead to fetch the next instructions. The pre-fetch stores the next two instructions and, on request, returns the stored instruction if the pre-fetched instruction matches the request.

For straight line code execution (that is, no branches), instructions are returned without any stalls. However, if there is a discontinuity in the program flow (such as `jump`, `call`, or wrapping of `do...loop`), the core is stalled for two clock cycles. The pre-fetch remembers the last discontinuity by storing the branch from and branch to address and the instruction after the branch. If the program flow comes back, it returns the correct instruction without stalling.

For example, if the Kalimba DSP enters into a `do...loop` that does not have any branches, the code executes without any stalls until it reaches the end of the `do...loop`. Then a discontinuity in the program flow occurs. The pre-fetch logic prediction of straight line code execution is then invalid because the code is nonsequential, which causes two stalls for the program counter to return to the start of the loop. However, the next time, the pre-fetch remembers the jump, and no stalls occurs.

4.10 Prefixes

A prefix instruction is required to use a 24-bit constant in a Type B instruction. The assembler **kalasm3** adds this automatically, as required. The top 8 bits of the constant are in the prefix instruction, with the remaining 16 bits in the main instruction.

For example:

```
r0 = r1 + 0x12345;
```

becomes

```
Prefix (0x1);
r0 = r1 + 0x2345;
```

Without a prefix instruction, the hardware must convert a 16-bit constant to 24 bits to be consistent with the contents of a register. Depending on the associated instruction, there are three ways this conversion can occur. Therefore, whether a prefix is needed is different for different instructions. The three types of conversion are as follows:

- Logical operations (that is, **AND**, **OR**, and **XOR**, see Section 6.4)). The top 8 bits are 0. The constant is placed in the bottom 16 bits.
- Fractional multiplies (see Section 6.8 and 6.9). For example:


```
r0 = r1 * 0.5 (frac); // 0.5 is 0x400000 as a 24-bit fractional
rMAC = r0 * 0x123400; // Fractional specified in hex format
rMAC = rMAC + r0 * 0x123400;
```

In these examples, the bottom 8 bits are 0. The constant is in the top 16 bits, and no prefix is required.
- For everything else, the constant is sign-extended. That is, the top 8 bits are the same as the top bit of the constant. The constant is in the bottom 16 bits.

NOTE Even where the constant is entered as a fractional number, this is the conversion method that is used. For example:

```
r0 = r1 + 0.5;
```

translates to:

```
r0 = r1 + 0x400000;
```

This requires a prefix, see [Table 4-6](#).

Table 4-6 Prefix Requirement Examples

Type	Requires Prefix	No Prefix Required
Logical	0x123456	0x003456
Fractional multiply	0x123456	0x123400
Everything else	0x123456 0x009456 0xff7456	0x003456 0x007456 0xff9456

4.11 Circular Buffers

A circular or ring buffer is a vector that uses a single, fixed-size buffer as if it were connected end to end. Some hardware assistance is required.

Hardware assistance is available only in indexed memory accesses (see Section 6.18), when the appropriate length register is set (see Section 4.7.3). Circular buffers can be of any length less than 64 K words.

Kalimba Architecture 5 contains a circular buffer mode, called *base register mode*. In this mode, any address can be a valid start address for a circular buffer.

It is possible to use circular buffers with or without base registers. The advantage of using base registers is that circular buffers can be at any location in the RAM. There are strict alignment requirements if base registers are not used. The advantage of not using base registers is a reduction in setup code required before each use.

4.11.1 Circular Buffers Without Use of Base Registers

Use one of the `DMCIRC`, `DM1CIRC`, or `DM2CIRC` directives to declare a circular buffer. This ensures that the buffer has a valid start address. Valid start addresses have zeros in all of the locations set in the offset mask.

The following is an example of how to use circular buffers:

```
.VAR/DMCIRC buf[13];    // say buf gets placed at 0x70
I0 = &buf;              // I0 = 0x70
L0 = LENGTH(buf);       // L0 = 0xd
r0 = M[I0,-1];          // Read from address 0x70
                        // (I0 = 0x7c after post-modify and wrap)
```

The following pseudocode explains the logic used to calculate increments of index registers. First, create an offset mask:

```
offset_mask = 2^(ceil(log2(length(buf)))) - 1
```

Then, calculate the new value of `I0` according to the following pseudocode, where the increment in the example equals -1:

```

base = I0 & ~offset_mask
offset = (I0 & offset_mask) + increment
if offset >= 0 && offset > length(buf)
    offset = offset - length(buf)
else if offset < 0
    offset = offset + length(buf)
I0 = base | offset

```

4.11.2 Circular Buffers With Use of Base Registers

Use one of the standard `DM`, `DM1`, or `DM2` directives to declare a circular buffer using base registers. The base register is set to the start address of the buffer.

An example of this mode is as follows:

```

.VAR/DM buf[13];           // say buf gets placed at 0x1234
I0 = &buf;                  // I0 = 0x1234
push I0;
pop B0;                     // B0 = 0x1234
L0 = LENGTH(buf);          // L0 = 0xd
r0 = M[I0,-1];              // Read from address 0x1234
                             // (I0 = 0x1240 after post-modify and wrap)

```

4.12 Zero Overhead Looping: do loop

For zero overhead looping, instructions between `do` and `loop` execute until register `r10` is 0. `r10` is decremented by 1 each loop. This decrement executes just before the last instruction in the loop. Use the value of `r10` inside a `do...loop` with caution.

If `r10` is 0 at the start of the loop, then no loop instructions execute and a jump to the loop-end label occurs.

If a `do...loop` is executed in an ISR, `r10` and the memory mapped registers `MM_DOLOOP_START` and `MM_DOLOOP_END` must be saved.

For example:

```

r0 = 1;
r1 = 1;
do loop;
    r2 = r0 + r1;           // this code will
    r0 = r1 + r2;           // be executed
    r1 = r2 + r0;           // 100 times
loop:

```

Kalimba Architecture 5 contains a 32-word `do...loop` cache to reduce power consumption while executing a `do...loop`. For many programs, a significant proportion of processing time is inside these loops. Therefore, this improvement leads to a considerable reduction in average power consumption.

The cache prevents the core from needing to fetch instructions from program RAM, thus requiring less power. Program RAM is already accessible in a single cycle, so there is no speed advantage. If code

is being run from flash or ROM, the do loop cache still provides a power-saving benefit over the 1024-word or 512-word flash or ROM cache in RAM.

4.13 Debug

The Kalimba DSP debugging hardware provides several features to an external debugger.

The features are:

- Reset, run, stop, step
- Setting and reading of the PC
- 8 program breakpoints
- 2 data memory breakpoints (read, write, or read and write), each covering an address range start to end
- Read and write of register values
- Read and write memory locations, as seen by the DSP on any of its three memory buses, PM, DM1, and DM2
- Read and clear profiling counters:
 - ☐ Number of clock cycles
 - ☐ Number of instructions executed
 - ☐ Number of stall cycles

5 Memory Organisation

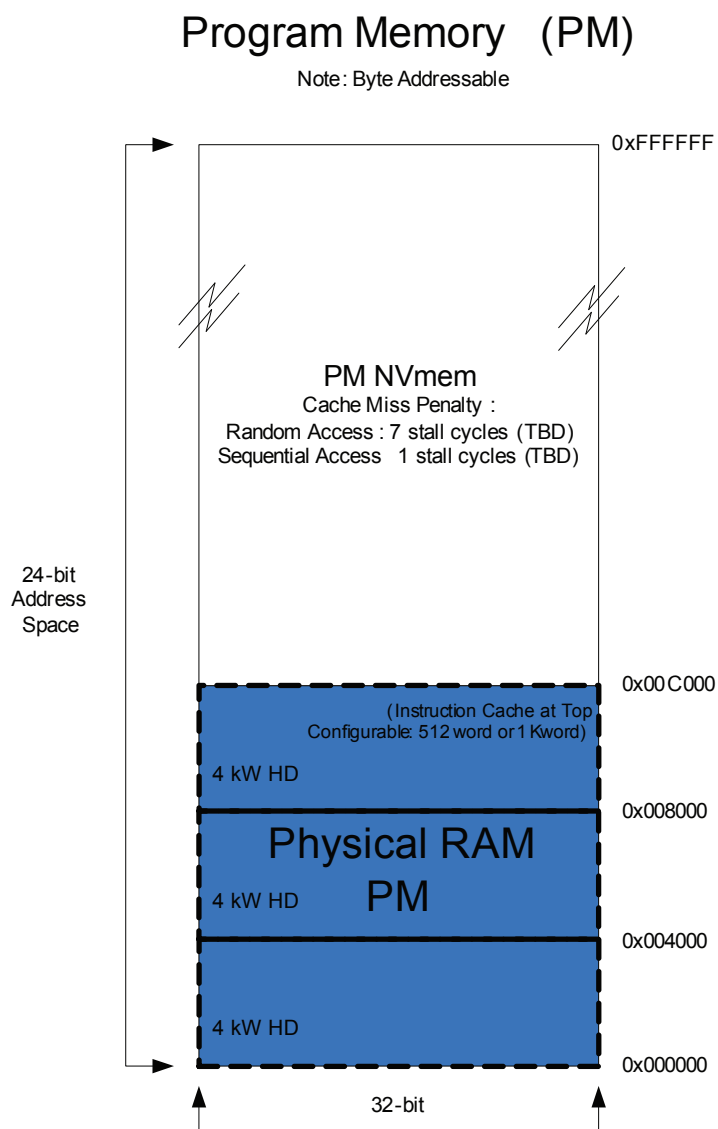
The Kalimba DSP core includes a single program memory, and dual data memory banks.

The Kalimba DSP core memory fetures include:

- A single 32-bit program memory, PM, with a 24-bit address space
- Two 24-bit data memory banks, DM1 and DM2, each with a 23-bit address space.

[Figure 5-1](#) shows the 32-bit PM. [Figure 5-2](#) shows the 24-bit data memory banks.

Accessing all three memories simultaneously is possible in the same clock cycle, assuming there are no conflicts. This is known as a *threebank Harvard architecture*. Conflicts introduce an appropriate number of wait cycles.



G-TW-0013732.1.3

Figure 5-1 Program Memory Organisation for CSR8675

CSR8675 device has the following physical RAM for the Kalimba DSP:

- DM1 = 32K x 24-bit
- DM2 = 32K x 24-bit
- PM = 12K x 32-bit

The BlueCore MCU initialises the Kalimba DSP. During initialisation, program and data coefficient download to the DSP occurs through auto-incrementing memory-mapped registers in the MCU. The

MCU then sets the initial clock frequency for the Kalimba DSP to use before starting it running. An application running on the MCU invokes the program download and Kalimba DSP initialisation.

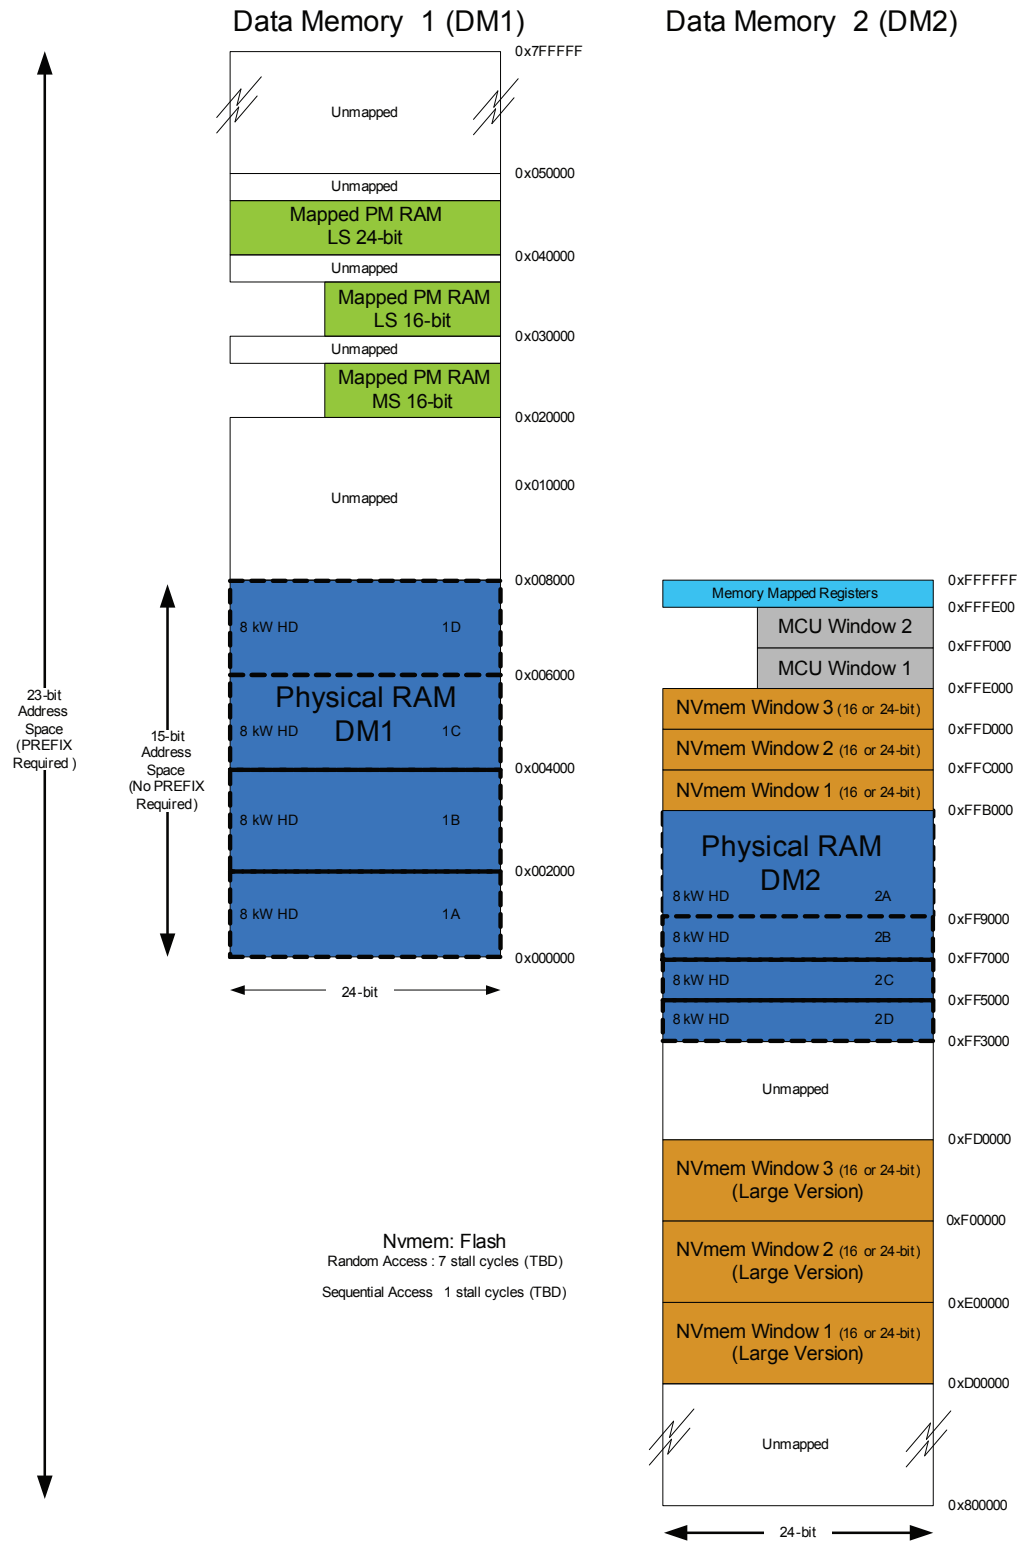


Figure 5-2 Data Memory Organisation for CSR8675

5.1 Memory Map

The memory organisation in Figure 5.1 and Figure 5.2 breaks down into individual memory maps described in Table 5.1, Table 5.2, and Table 5.3. The tables show the DSP memory layout for QTIL devices containing Kalimba Architecture 5, using CSRC3707 as an example.

5.1.1 PM Memory Map

The program memory map for Kalimba contains the physical program memory RAM.

Some QTIL devices contain a low-power RAM region, which is available as cache for the PM flash or ROM space. The remaining 16 M words can map into flash or ROM. Memory-mapped registers in the DSP select the flash or ROM start address and size of the region of flash or ROM mapped. The MCU can control the priority given to flash or ROM accesses between the Kalimba DSP and MCU.

Table 5-1 PM Memory Map

Label	Description
Physical RAM PM	General-purpose RAM for program code
Physical RAM PM (instruction cache)	Cache for flash or ROM PM space, or general-purpose RAM if flash or ROM is not used Consists of low-power memory in some CSR devices
PM NVMem	Mapped to flash or ROM

5.1.2 DM1 Memory Map

The memory map of the first data memory bank DM1 contains data memory for the DSP, some of which is low-power data RAM in some CSR devices and a window of program memory split into the most and least significant halves, and including a 24-bit window. The remaining areas of the memory map allow for future variants.

Table 5-2 DM1 Memory Map

Label	Description
Physical RAM DM1	General purpose RAM (data memory 1)
Unmapped	Available for RAM expansion in future variants
Mapped PM RAM (MS 16 bits)	Mapped to program memory MS 16 bits (equal in size to program memory)
Mapped PM RAM (LS 16 bits)	Mapped to program memory LS 16 bits (equal in size to program memory)
Mapped PM RAM (LS 24 bits)	Mapped to program memory LS 24 bits (equal in size to program memory)

5.1.3 DM2 Memory Map

The memory map of the second data memory bank, DM2, contains:

- General purpose data RAM for the DSP, a portion of which is low-power data RAM in some QTIL devices
- Three 1 M words or 4 K words windows into the flash or ROM, which is available to the DSP for items such as constant coefficient tables
- Two MCU windows (4 K and 3.5 K words) into the MCU memory, which enable control information and message passing
- 512 words, which are reserved for the memory-mapped I/O for the DSP

Table 5-3 DM2 Memory Map

Label	Description
Unmapped	Available for RAM expansion in future variants
NVmem window 1 (16-bit or 24-bit) (large version)	Large, nonvolatile memory window 1 (1 M word)
NVmem window 2 (16-bit or 24-bit) (large version)	Large, nonvolatile memory window 2 (1 M word)
NVmem window 3 (16-bit or 24-bit) (large version)	Large, nonvolatile memory window 3 (832 K word)
Physical RAM DM2	General-purpose RAM (data memory 2)
NVmem Window 1 (16-bit or 24-bit)	Nonvolatile memory window 1 (4 K words)
NVmem Window 2 (16-bit or 24-bit)	Nonvolatile memory window 2 (4 K words)
NVmem window 3 (16-bit or 24-bit)	Nonvolatile memory window 3 (4 K words)
MCU window 1	Window into MCU memory 1 (4 K words)
MCU window 2	Window into MCU memory 2 (3.5 K words)
Memory-mapped registers	DSP memory-mapped I/O registers (512 words)

6 Instruction Set

The instruction set for the Kalimba DSP is suited to an algebraic assembler, rather than the mnemonic assemblers found in traditional MCUs. Algebraic assemblers suit the architecture of a DSP because they can express complex and parallel instructions in an understandable way. The Kalimba DSP architecture is designed to execute single-cycle instructions. There are exceptions (for example, the divide instruction).

Table 6-1 shows the instruction set syntax described using notation conventions.

Table 6-1 Instruction Set Descriptions

Parallel lines	Vertical parallel bars enclose lists of syntax options. Selection of one of the options is required.
Angled brackets <non-bold italics>	Anything in non-bold italics enclosed by angled brackets is an optional part of the instruction statement.
A, B, C	Denotes a register operand. By default, the register is selected from the list of Bank 1 registers. If subscripted with <i>Bank 1/2</i> , then the register can be chosen from either Bank 1 or Bank 2. If subscripted with <i>Bank 1/2/3</i> , then the register can be chosen from Bank 1, Bank 2, or Bank 3.
k_{16} , k_7 , k_n	Denotes a constant, where the subscripted number indicates the size of the number in bits.
M[x]	Data in memory location with address x.
M[i,m]	<i>i</i> is the index register. <i>m</i> is the modify register.
cond	A condition code (for example, NZ).
MEM_ACCESS_1 MEM_ACCESS_2	Represents a memory access instruction that can be appended to an instruction where indicated.
rMAC(B)	Represents either rMAC or rMACB registers.

6.1 ADD and ADD with CARRY

Syntax

Type A:	$\langle \text{if cond} \rangle$ $C = A + B$ $C = A + M[B]$ $C = M[A] + B$ $M[C] = A + B$	$\langle +\text{Carry} \rangle$ $\langle \text{MEM_ACCESS_1} \rangle ;$
Example:	if Z r3 = r1 + M[r2] + Carry, r4 = M[I0, M0];	
Type B:	$C = A + k_{16}$ $C = A + M[k_{16}]$ $C = M[A] + k_{16}$ $M[k_{16}] = A + C$	$\langle +\text{Carry} \rangle ;$
Example:	r3 = M[r1] + 10 + Carry;	
Type C:	$C = C + A$ $C = C + M[A]$	$\langle +\text{Carry} \rangle$ $\langle \text{MEM_ACCESS_1} \rangle$ $\langle \text{MEM_ACCESS_2} \rangle ;$
Example:	r3 = r3 + M[r1] + Carry, r4 = M[I0,M0], r5 = M[I4,M1];	

Description:

Test the optional condition and, if TRUE, perform the addition. If the condition is FALSE, perform a NOP but *MEM_ACCESS_1* is still carried out. Omitting the condition performs the addition unconditionally. The addition operation adds the first source operand to the second source operand and, if designated by the *+ Carry* notation, adds the ALU carry bit, C. The result is stored in the destination operand. The operands may be either one of the 16 Bank 1 registers, a 16-bit sign extended constant (24-bit with prefix instruction), or memory pointed to by a register or a constant.

Flags Generated:

Z	Set if the result equals 0 and cleared otherwise	N	Set if the result is negative and cleared otherwise
V	Set if an arithmetic overflow occurs and cleared otherwise	C	Set if a carry is generated and cleared otherwise

NOTE If one of the source operands is Null then a LOAD/STORE is assumed. Therefore, the C and V flags are cleared.

If all operands are Null then a NOP is assumed. Therefore, all flags are unchanged.

When writing to rFlags the result is stored in rFlags and so the flags above are not applicable.

Additions and subtractions can saturate on overflow by setting the ARITHMETIC_MODE memory-mapped register. This causes saturation of signed numbers (positive saturation to 0x7fffffff, negative saturation to 0x8000000).

kalasm3 can understand the following instructions because the 0 can be implemented using the Null register.

```
if Z r1 = 0, r4 = M[I1, M1];
```

6.2 SUBTRACT and SUBTRACT With Borrow

Syntax:

Type A:	$\langle \text{if cond} \rangle$ $\begin{array}{l} C = A - B \\ C = A - M[B] \\ C = M[A] - B \\ M[C] = A - B \end{array}$	$\langle \text{-Borrow} \rangle$ $\langle \text{MEM_ACCESS_1} \rangle;$
Example:	<pre>if Z r3 = r1 - r2 - Borrow, r4 = M[I0, M0];</pre>	
Type B:	$\begin{array}{l} C = A - k_{16} \\ C = A - M[k_{16}] \\ C = M[A] - k_{16} \\ M[k_{16}] = A - C \end{array}$	$\langle \text{-Borrow} \rangle$
Example:	<pre>r3 = M[r1] - 10;</pre>	
Type C:	$\begin{array}{l} C = C - A \\ C = C - M[A] \end{array}$	$\langle \text{-Borrow} \rangle$ $\langle \text{MEM_ACCESS_1} \rangle$ $\langle \text{MEM_ACCESS_2} \rangle;$
Example:	<pre>r3 = r3 - r1 - Borrow, r4 = M[I0, M0], r5 = M[I4, M1];</pre>	

Description

Test the optional condition and, if **TRUE**, perform the subtraction. If the condition is **FALSE**, perform a **NOP** but **MEM_ACCESS_1** is still carried out. Omitting the condition performs the subtraction unconditionally. The subtraction operation subtracts the second source operand from the first source operand and optionally, if designated by the **- Borrow** notation, subtracts the ALU carry bit, **B**. The result is stored in the destination operand. The operands may be either one of the 16 Bank 1 registers, a 16-bit sign extended constant (24-bit with prefix instruction), or memory pointed to by the register or a constant.

Flags Generated:

Z	Set if the result equals 0 and cleared otherwise	N	Set if the result is negative and cleared otherwise
V	Set if an arithmetic overflow occurs and cleared otherwise	C	Cleared if a borrow is generated and set otherwise

NOTE If one of the source operands is **Null** then a negate is assumed. Therefore, the **C** and **V** flags are cleared as appropriate.

When writing to **rFlags** the result is stored in **rFlags** and so the flags above are not applicable.

Additions and subtractions can saturate on overflow by setting the **ARITHMETIC_MODE** memory-mapped register. This causes saturation of signed numbers (positive saturation to **0x7fffff**, negative saturation to **0x800000**).

6.3 Bank 1/2 Register Operations: ADD and SUBTRACT

Syntax:

Type A:	$\begin{array}{l} \text{C}_{\text{BANK } 1/2} = \text{A}_{\text{BANK } 1/2} + \text{B}_{\text{BANK } 1/2} \\ \text{C}_{\text{BANK } 1/2} = \text{A}_{\text{BANK } 1/2} - \text{B}_{\text{BANK } 1/2} \end{array}$	$\langle \text{MEM_ACCESS_1} \rangle;$
Example:	<pre>if Z I0 = I4 + r2, r1 = M[I1,M1];</pre>	
Type B:	$\begin{array}{l} \text{C}_{\text{BANK } 1/2} = \text{A}_{\text{BANK } 1/2} + k_{16} \\ \text{C}_{\text{BANK } 1/2} = \text{A}_{\text{BANK } 1/2} - k_{16} \\ \text{C}_{\text{BANK } 1/2} = k_{16} - \text{A}_{\text{BANK } 1/2} \end{array}$;
Example:	<pre>I0 = r2 + 5;</pre>	
Type C:	$\begin{array}{l} \text{C}_{\text{BANK } 1/2} = \text{C}_{\text{BANK } 1/2} + \text{A}_{\text{BANK } 1/2} \\ \text{C}_{\text{BANK } 1/2} = \text{C}_{\text{BANK } 1/2} - \text{A}_{\text{BANK } 1/2} \end{array}$	$\langle \text{MEM_ACCESS_1} \rangle \quad \langle \text{MEM_ACCESS_2} \rangle;$
Example:	<pre>r2 = r2 + I2, r0 = M[I0,M0], r1 = M[I4,M1];</pre>	

Description:

Test the optional condition and, if **TRUE**, perform the specified cross-bank addition or subtraction. If the condition is **FALSE**, perform a NOP, but **MEM_ACCESS_1** is still carried out. Omitting the condition performs the addition or subtraction unconditionally. The operands may be either one of the 16 Bank 1 or 16 Bank 2 registers or a 16-bit sign extended constant (24-bit with prefix instruction).

Flags Generated:

Z	Set if the result equals 0 and cleared otherwise	N	Set if the result is negative and cleared otherwise
V	Set if an arithmetic overflow occurs and cleared otherwise	C	For addition, set if a carry is generated For subtraction, cleared if a borrow is generated

NOTE If one of the source operands is **Null**, then a **LOAD/STORE** is assumed. The **C** and **V** flags are cleared.

When writing to **rFlags**, the result is stored in **rFlags**. The flags are not applicable.

Additions and subtractions can saturate on overflow by setting the **ARITHMETIC_MODE** memory-mapped register. This causes saturation of signed numbers (positive saturation to **0x7fffffff**, negative saturation to **0x80000000**). This only occurs where the destination register is a Bank 1 24-bit register.

6.4 Logical Operations: AND, OR and XOR

Syntax:

Type A:	<div> <div><if cond></div> <div> C = A AND B C = A OR B C = A XOR B </div> </div>	<MEM_ACCESS_1>;
Example:	<pre>if Z r3 = r1 AND r2, r0 = M[I0,M0];</pre>	
Type B:	<div> C = A AND k₁₆ C = A OR k₁₆ C = A XOR k₁₆ </div>	;
Example:	<pre>r3 = r1 XOR 10;</pre>	
Type C:	<div> C = C AND A C = C OR A C = C XOR A </div>	<MEM_ACCESS_1> <MEM_ACCESS_2>;
Example:	<pre>r3 = r3 OR r1, r0 = M[I0,M0], r2 = M[I4,M1];</pre>	

Description:

Test the optional condition and, if **TRUE**, perform the specified bit-wise logical operation (logical **AND**, **OR**, or **XOR**). If the condition is **FALSE**, perform a **NOP**, but **MEM_ACCESS_1** is still carried out. Omitting the condition performs the operation unconditionally. The operands can be one of the 16 Bank 1 registers or a 16-bit, zero-padded at the most significant end, constant (24-bit with prefix instruction).

Flags Generated:

Z	Set if the result equals zero and cleared otherwise	N	Set if the result is negative and cleared otherwise
V	Cleared	C	Cleared

6.5 Shifter: LSHIFT and ASHIFT

Syntax:

Type A:	<code><if cond> C = A OP B</code>	<code><MEM_ACCESS_1>;</code>
Example:	<pre>if Z r3 = r2 LSHIFT r1, r0 = M[I0,M0];</pre>	
Type B:	$C = A \text{ OP } k_7$ $rMAC(B)0 = A \text{ OP } k_7$ $rMAC(B)12 = A \text{ OP } k_7$ $rMAC(B)2 = A \text{ OP } k_7$ $rMAC(B) = A \text{ OP } k_7 \text{ (LO)}$ $rMAC(B) = A \text{ OP } k_7 \text{ <(MI)>}$ $rMAC(B) = A \text{ OP } k_7 \text{ (HI)}$ $C = k_{16} \text{ OP } A$	<code>;</code>
Example:	<code>rMAC0 = r2 LSHIFT 4;</code>	
Type C:	<code>C = C OP A</code>	<code><MEM_ACCESS_1> <MEM_ACCESS_2>;</code>
Example:	<pre>r2 = r2 ASHIFT r7, r5 = M[I0,M2], r1 = M[I4,M1];</pre>	

Description:

Test the optional condition and, if **TRUE**, perform the specified shift operation (arithmetic or logical). If the condition is **FALSE**, perform a **NOP**, but **MEM_ACCESS_1** is still carried out. Omitting the condition performs the shift unconditionally. A positive number causes shifting to the left. A negative number causes shifting to the right. For an arithmetic shift to the right, sign extension bits are added, as needed. If overflow occurs in an arithmetic shift (that is, nonsign bits being shifted out), then the overflow flag is set. The result is saturated to $2^{23}-1$ or -2^{23} , depending on the sign of the input. No rounding occurs for **ASHIFT** or **LSHIFT**. The operands can be any one of the 16 Bank 1 registers or a constant specified in the instruction.

Flags Generated:

Z	Set if the result equals 0 and cleared otherwise	N	Set if the result is negative, and cleared otherwise
V	ASHIFT: Set if an arithmetic overflow occurs, and cleared otherwise LSHIFT: Cleared	C	Cleared

NOTE OP is either **ASHIFT** (arithmetic) or **LSHIFT** (logical).

If **rMAC** or **rMACB** is the source operand, the full 56 bits are used as input to the shifter. The output of the shifter is 24-bit in this instruction.

rFlags cannot be a destination operand for **ASHIFT** and **LSHIFT**.

For Type B instructions, the destination operand can be a Bank 1 register or:

- **rMAC (B) 0, rMAC (B) 12, or rMAC (B) 2, causing the other bits of rMAC (B) to be unaffected (writing to rMAC (B) 12 writes the data into rMAC (B) 1 and causes sign extension (ASHIFT) or zero fill (LSHIFT) into rMAC (B) 2)**
- **rMAC (B) , with a data format tag (LO, MI, HI) to select which word of rMAC (B) the 24-bit result from the shifter should be written to. The other bits of rMAC (B) are sign-extended or zero-padded, as appropriate (see Section 7.13 and Section 7.14). MI is the default tag if none is supplied.**

The source register also can be a 16-bit, sign-extended constant (or 24-bit constant with prefix).

6.6 LSHIFT and ASHIFT (56 bit)

Syntax:

Type A:	<code><if cond> rMAC(B) = A OP B (56 bit) <MEM_ACCESS_1>;</code>
Example:	<code>if Z rMACB = r2 LSHIFT r1 (56 bit), r0 = M[I0,M0];</code>
Type B:	<code>rMAC(B) = A OP k₇ (56 bit) ;</code>
Example:	<code>rMAC = r2 LSHIFT 4 (56 bit); rMACB = 0x1234 ASHIFT r3 (56 bit);</code>
Type C:	<code>rMAC(B) = rMAC(B) OP A (56 bit) <MEM_ACCESS_1> <MEM_ACCESS_2>;</code>
Example:	<code>rMACB = rMACB ASHIFT r7 (56 bit), r5 = M[I0,M2], r1 = M[I4,M1];</code>

Description:

Test the optional condition and, if **TRUE**, perform the specified shift operation (arithmetic or logical). If the condition is **FALSE**, perform a **NOP**, but **MEM_ACCESS_1** is still carried out. Omitting the condition performs the shift unconditionally. A positive number causes shifting to the left. A negative number causes shifting to the right. For an arithmetic shift to the right, sign extension bits are added as needed. If overflow occurs in an arithmetic shift, that is, nonsign bits being shifted out, then the overflow flag is set. The result is saturated to $2^{47}-1$ or -2^{47} , depending on the sign of the input. No rounding occurs for **ASHIFT** or **LSHIFT**. The operands can be one of the 16 Bank 1 registers or a constant specified in the instruction.

Flags Generated:

Z	Set if the result equals zero and cleared otherwise	N	Set if the result is negative and cleared otherwise
V	ASHIFT: Set if an arithmetic overflow occurs and cleared otherwise LSHIFT: Cleared	C	Cleared

NOTE OP is either **ASHIFT** (arithmetic) or **LSHIFT** (logical).

If **rMAC** or **rMACB** is the source operand, the full 56 bits are used as input to the shifter. The output of the shifter is always 56-bit in this instruction.

For Type B instructions, the source register can be a 16-bit, sign-extended constant (or 24-bit constant with prefix).

6.7 rMAC(B) Move Operations

Syntax:

Type B:	rMAC(B)0	=	rMAC(B)0	;
	rMAC(B)2		rMAC(B)1	
			rMAC(B)2	
			A	

Example: `rMAC0 = rMAC1;`

rMAC(B)12	=	rMAC(B)0	(SE)	;
		rMAC(B)1	(ZP)	
		rMAC(B)2		
		A		

Example: `rMAC12 = rMAC0 (SE);`

C	=	rMAC(B)0	;
		rMAC(B)1	

Example: `r3 = rMAC0;`

C	=	rMAC(B)2	(SE)	;
			(ZP)	

Example: `r3 = rMAC2 (ZP);`

Description:

These are move instructions, implemented as a special case of `LSHIFT` and `AShift`, to support loading and reading of the individual sections of the `rMAC (B)` registers (`rMAC (B) 2`, `rMAC (B) 1`, and `rMAC (B) 0`). See Figure 4-2.

When writing to `rMAC (B) 1`, the data is either sign-extended or zero-padded into `rMAC (B) 2`. The format specifiers, `SE` (sign extend) and `ZP` (zero pad), are required to specify how `rMAC (B) 2` is filled.

Flags Generated:

Z	Set if the result equals zero and cleared otherwise	N	Set if the result is negative and cleared otherwise
V	Cleared	C	Cleared

NOTE These operations are always a Type B instruction, that is, no parallel memory reads are performed.

6.8 Multiply: Signed 24-bit Fractional and Integer

Syntax:

Type A: `<if cond> C = A * B` $\left| \begin{array}{l} \text{(frac)} \\ \text{(int) } \text{<(sat)>} \end{array} \right|$ `<MEM_ACCESS_1>;`

Example: `if Z r3 = r2 * r1 (int) (sat),
 r5 = M[I0,M2];`

Type B: `C = A * k16` $\left| \begin{array}{l} \text{(frac)} \\ \text{(int) } \text{<(sat)>} \end{array} \right|$ `;`

Example: `r6 = r2 * 0.34375 (frac);`

Type C: `C = C * A` $\left| \begin{array}{l} \text{(frac)} \\ \text{(int) } \text{<(sat)>} \end{array} \right|$ `<MEM_ACCESS_1> <MEM_ACCESS_2>;`

Example: `r2 = r2 * r7 (int) (sat),
 r0 = M[I0,1],
 r1 = M[I4,-1];`

Description:

Test the optional condition and, if **TRUE**, perform the specified multiply operation (fractional or integer). If the condition is **FALSE**, perform a **NOP**, but **MEM_ACCESS_1** is still carried out. Omitting the condition performs the operation unconditionally. A fractional multiply, *(frac)*, treats the source and destination operands as fractional numbers with 2^{23} representing +1.0 and -2^{23} representing -1.0. An integer multiply, *(int)*, treats the source and destination operands as integer numbers. Optionally, if designated by the *(sat)* notation, the result of an integer multiply saturates if overflow occurs. Unbiased rounding is always done for a fractional multiply operation. The operands may be either one of the 16 Bank 1 registers or a 16-bit constant (24-bit with prefix instruction). See Appendix A on number representation for information on multiply operations.

Flags Generated:

Z	Set if the result equals zero and cleared otherwise	N	Set if the result is negative and cleared otherwise
V	frac: Cleared int: Set if an arithmetic signed overflow occurs and cleared otherwise	C	Cleared

NOTE A saturated fractional multiplication has no significance therefore *(sat)* is not an option with *(frac)*.

The 16-bit constant has a different meaning for *(int)* and *(frac)* operations:

- For *(frac)*, the constant is left justified to 24 bits by adding 8 zeros as the LSBs. This enables the 16-bit constant to represent fixed-point fractional numbers between +1.0 and -1.0.
- For *(int)*, the constant is sign extended to 24 bits by adding 8 ones or zeros as the MSBs. This enables the 16-bit constant to represent numbers between +32767 and -32768.

Unbiased rounding of a fractional multiply is carried out according to the following pseudocode:

```
temp = (operand1 * operand2) << 1;    // fractional multiply
result = (temp + 0x800000) >> 24;      // round result
if ((temp & 0xFFFFFFF) == 0x800000)    // unbias above rounding
    result = result & 0xFFFFFFE;
```

When the number to be rounded is precisely half way between the 2 nearest values of the result, odd values are rounded away from zero and even values towards zero, yielding a zero large sample bias assuming uniformly distributed values.

6.9 MULTIPLY and ACCUMULATE (56 bit)

Syntax:

Type A:	<i><if cond></i>	$C'' = A * B$	<i><(SS)></i>	<i><MEM_ACCESS_1>;</i>
		$C'' = C'' + A * B$	<i><(SU)></i>	
		$C'' = C'' - A * B$	<i><(US)></i>	
			<i><(UU)></i>	

Example: `if Z C'' = C'' + r1*r2 (SS),
r5 = M[I0,M0];`

Type B:	$C'' = A * k_{16}$	<i><(SS)></i>	;
	$C'' = C'' + A * k_{16}$	<i><(SU)></i>	
	$C'' = C'' - A * k_{16}$	<i><(US)></i>	
		<i><(UU)></i>	

Example: `rMAC = rMAC + r1 * 0.24254 (SS);`

Type C:	$rMAC = C * A$	<i><(SS)></i>	<i><MEM_ACCESS_1></i>	<i><MEM_ACCESS_2>;</i>
	$rMAC = rMAC + C * A$	<i><(SU)></i>		
	$rMAC = rMAC - C * A$	<i><(US)></i>		
		<i><(UU)></i>		

Example: `rMAC = rMAC - r3 * r1 (SS),
r2 = M[I0,1],
r1 = M[I4,-1];`

Description:

Test the optional condition and, if **TRUE**, perform the specified multiply/accumulate. If the condition is **FALSE**, perform a **NOP**, but *MEM_ACCESS_1* is still carried out. Omitting the condition performs the operation unconditionally. The data format field to the right of the operands specifies whether each respective operand is in signed (**S**) or unsigned (**U**) format. The effective binary point is between bits 47 and 46. The operands can be one of the 16 Bank 1 registers or a 16-bit, left-justified constant (24-bit with prefix instruction). See Appendix A for information about multiply operations.

Flags Generated:

Z	Set if the result equals 0 and cleared otherwise	N	Set if the result is negative and cleared otherwise
V	Set if overflow occurs past the 56 th bit and cleared otherwise	C	Cleared

NOTE Where *(SS)* is the default if no data format is specified.

The 16-bit constant is left justified to 24 bits by adding 8 zeros as the LSBs. This enables the 16-bit constant to represent fixed-point fractional numbers between +1.0 and -1.0.

To get the result of the equivalent integer multiplication, shift the result to the right by 1 bit.

For Type A and B instructions not using the supplementary ADD/SUB operation, the accumulator register C" may be selected from `rMAC`, `rMACB`, `r0`, `r1`, `r2`, `r3`, `r4` or `r5`, and the data format of the operators may be chosen to be signed (S) or unsigned (U).

6.10 MULTIPLY and ACCUMULATE (56 bit) with ADD/SUB

Syntax:

Type A:

<if cond>	$C'' = A * B$	$\langle r0 = D + E \rangle /$	$\langle MEM_ACCESS_1 \rangle;$
	$C'' = C'' + A * B$	$\langle SS \rangle / \langle SU \rangle /$	
	$C'' = C'' - A * B$	$\langle US \rangle / \langle UU \rangle$	

Example:

```
if NZ r1 = r1 + r7 * r8, r0 = r1 - rMACB,
    r2 = M[I1,M1];

if C rMACB = rMACB + r4 * r5, r0 = r2 + rMACB,
    M[I2,M2] = r1;

if Z r4 = r4 + r1*r2 (SU),
    r5 = M[I0,M0];
```

Type B:

$C'' = A * k_{16}$	$\langle r0 = D + E \rangle /$;
$C'' = C'' + A * k_{16}$	$\langle SS \rangle / \langle SU \rangle /$	
$C'' = C'' - A * k_{16}$	$\langle US \rangle / \langle UU \rangle$	

Example:

```
rMAC = rMAC + r1 * 0.24254 (SS);
```

Type C:

$rMAC(B) = C' * A$	$r0 = D +/-$	$\langle MEM_ACCESS_1 \rangle \quad \langle MEM_ACCESS_2 \rangle;$
$rMAC(B) = rMAC(B) + C' * A$	$rMAC(B)$	
$rMAC(B) = rMAC(B) - C' * A$		

Example:

```
rMAC = rMAC - r3 * r1, r0 = r1 + rMACB,
r2 = M[I0,1], r1 = M[I4,-1];
```

Description:

Test the optional condition and, if **TRUE**, perform the specified multiply/accumulate along with the separate add/subtract operation. If the condition is **FALSE**, perform a **NOP**, but **MEM_ACCESS_1** is still carried out. Omitting the condition performs the operation unconditionally. The effective binary point is between bits 47 and 46. See [Number Representation](#) for information on multiply operations. See [Type B Stack Encoding](#) for tables describing all possible combinations of this instruction.

Flags Generated:

Z	Set if the result equals zero and cleared otherwise	N	Set if the result is negative and cleared otherwise
V	Set if overflow occurs past the 56 th bit and cleared otherwise	C	Cleared

NOTE To get the result of the equivalent integer multiplication, shift the result to the right by 1 bit.

For the supplementary **ADD/SUB** operation, Reg D can be selected from **r1** and **r2**.

For Type A and B instructions using the supplementary **ADD/SUB** operation, Reg E can be selected from **rMAC** or **rMACB**.

The accumulator register C'' may be selected from rMAC, rMACB, r1 or r2, and the data format of both operators is assumed to be signed (SS).

For Type C instructions, the second operand of the add/sub is always an rMAC (B) register, and is the opposite rMAC (B) register to that used for the multiply part. Reg C' specifies the first multiplicand, and may select from: rMACB, rMAC, r0, r1, r2, r3, r4 or r5. The data format of both operators is assumed to be signed (SS).

Flags are set because of the result of the multiply operation; no flags are altered due to the result of the additional ADD/SUB operation.

6.11 LOAD/STORE with Memory Offset

Syntax:

Type A:	$\langle if\ cond \rangle$	$C = M[A + B]$ $M[A + B] = C$	$\langle MEM_ACCESS_1 \rangle;$
Example:		$if\ Z\ r3 = M[r1 + r2],$ $r4 = M[I0, M0];$	
Type B:		$C = M[A + k_{16}]$ $M[A + k_{16}] = C$;
Example:		$M[r3 + 6] = r1;$	
Type C:		$C = M[C + A]$	$\langle MEM_ACCESS_1 \rangle\ \langle MEM_ACCESS_2 \rangle;$
Example:		$r3 = M[r3 + r2],$ $r4 = M[I0, 1],$ $r5 = M[I4, -1];$	

Description:

Test the optional condition and, if **TRUE**, perform the specified **LOAD/STORE**, including memory offset. If the condition is **FALSE**, perform a **NOP**, but **MEM_ACCESS_1** is still carried out. Omitting the condition performs the **LOAD/STORE** unconditionally. The operands can be one of the 16 Bank 1 registers or a 16-bit constant specified in the instruction.

Flags Generated:

Z	Set if the result operand equals 0 and cleared otherwise	N	Set if the result is negative and cleared otherwise
V	Cleared	C	Cleared

NOTE Kalimba has two data memory banks, so the three memory accesses take at least two cycles, depending on which banks the addresses are in.

6.12 Extended LOAD/STORE with Memory Offset

Syntax:

Type A: $\langle \text{if cond} \rangle \left| \begin{array}{l} C_{1/2} = \text{Mxx}[A_{1/2} \pm B_{1/2}] \\ \text{Myy}[A_{1/2} \pm B_{1/2}] = C_{1/2} \end{array} \right| \langle \text{MEM_ACCESS_1} \rangle;$

Example: $\text{if } Z \text{ } r3 = M[r1 + r2],$
 $r4 = M[I0, M0];$

Type B: $\left| \begin{array}{l} C_{1/2} = \text{Mxx}[A_{1/2} \pm k_{11}] \\ \text{Myy}[A_{1/2} \pm k_{11}] = C_{1/2} \end{array} \right| ;$

Example: $M[r3 + 6] = r1;$

Description:

Test the optional condition and, if **TRUE**, perform the specified extended **LOAD/STORE**, including memory offset. If the condition is **FALSE**, perform a **NOP**, but **MEM_ACCESS_1** is still carried out. Omitting the condition performs the extended **LOAD/STORE** unconditionally. The operands can be one of the 16 Bank 1 registers, 16 Bank 2 registers, or an 11-bit constant specified in the instruction.

NOTE The Kalimba Architecture 5 DSP does not support subword access, so:

- Mxx is the read-a-word (32-bit) value
- Myy is the write-a-word (32-bit) value

Flags Generated:

Z	Set if the result operand equals zero and cleared otherwise	N	Set if the result is negative and cleared otherwise
V	Cleared	C	Cleared

NOTE Kalimba has two data memory banks, so the three memory accesses take at least two cycles, depending on which banks the addresses are in.

6.13 Sign Bits Detect and Block Sign Bits Detect

Syntax:

Type A:	<code><if cond> C = SIGNDET A <MEM_ACCESS_1>;</code>
Example:	<code>if Z r3 = SIGNDET rMAC, r4 = M[I0,M0];</code>
Type C:	<code>C = BLKSIGNDET A <MEM_ACCESS_1> <MEM_ACCESS_2>;</code>
Example:	<code>r3 = BLKSIGNDET r1, r4 = M[I0,1], r5 = M[I4,-1];</code>

Description:

SIGNDET returns the number of redundant sign bits of the source operand. For example:

0000 1101 0101 0101 1100 1111	has 3 redundant sign bits
1001 0101 0101 0100 0111 1111	has 0 redundant sign bits
0000 0000 0000 0000 0000 0001	has 22 redundant sign bits
1111 1111 1111 1111 1111 1111	has 23 redundant sign bits
0000 0000 0000 0000 0000 0000	has 23 redundant sign bits (special case)

If the condition is **FALSE**, perform a **NOP**, but **MEM_ACCESS_1** is still carried out. Omitting the condition performs the operation unconditionally.

Valid results are 0 to 23 for the 24-bit registers, and -8 to 47 for **rMAC (B)**.

BLKSIGNDET returns the smaller of the result of **SIGNDET** and the present value of the destination register. When performed on a series of numbers, it can derive the effective exponent of the number largest in magnitude.

Flags Generated:

Z	Set if the result equals 0 and cleared otherwise	N	Set if the result is negative and cleared otherwise
V	Cleared	C	Cleared

6.14 Divide Instruction

Syntax:

Type B:

Div = C/A	;
C = DivResult	
C = DivRemainder	

Example:

```
Div = C / r1;
r2 = DivResult;
r3 = DivRemainder;
```

Description:

The `Div = C/A` instruction initiates the divide block's multicycle 48-bit/24-bit integer divide. The overflow flag is set if the `DivResult` is wider than 24-bits. In this case, the result is saturated to -2^{23} or $2^{23} - 1$ and the remainder is invalid. There must be at least 12 cycles between the divide instruction and using the result or remainder. If the result or remainder request is before the 12 cycles have elapsed then program flow stalls until the result is ready. To carry out a fractional divide the value in `rMAC/rMAC (B)` first needs to be right-shifted by 1 bit before carrying out the divide operation. The result of the divide can also be pushed directly onto the stack, which stalls the DSP if the divide is not yet completed.

NOTE The value of `DivResult` and `DivRemainder` are available as memory-mapped registers called `MM_QUOTIENT` and `MM_REM`, respectively.

`DivResult` returns the result of the division truncated to an integer.

`DivRemainder` returns the remainder such that the equations below hold.

$\text{DivResult} = \text{trunc}(\text{Dividend} / \text{Divisor})$

$\text{DivRemainder} = \text{Dividend} - (\text{Divisor} \times \text{DivResult})$

Example:

Dividend/Divisor	DivResult	DivRemainder
5/3	1	2
-5/3	-1	-2
5/-3	-1	2
-5/-3	1	-2

Integer Divide example (86420/7 = 12345 remainder 5):

```
r0 = 86420;
rMAC = r0 ASHIFT 0 (LO);    // LS word of rMAC now equals 86420, with
                             // sign extension in higher bits;

r0 = 7;
Div = rMAC / r0;
r1 = DivResult;              // r1 = 12345
r2 = DivRemainder;           // r2 = 5
```


Fractional Divide example (0.25/0.75 = 0.3333):

```

rMACB = 0.25;
r0 = 0.75;
rMACB = rMACB ASHIFT -1;
Div = rMACB / r0;
r1 = DivResult;           // r1 = 0.33333;

```

Flags Generated:

After: Div = rMAC(B) / A:

Z	Unchanged	N	Unchanged
V	Set if a divide exception occurs (divide by 0 or overflow in DivResult) and cleared otherwise	C	Cleared

After: C = DivResult or C = DivRemainder:

Z	Set if the result equals 0 and cleared otherwise	N	Set if the result is negative and cleared otherwise
V	Cleared	C	Cleared

6.15 Stack Instructions

Syntax:

Type A:	$\langle \text{if cond} \rangle$ <div> push $C_{\text{BANK } 1/2/3}$ pop $C_{\text{BANK } 1/2/3}$ $C_{\text{BANK } 1/2/3} = M[\text{FP} + A]$ $M[\text{FP} + A] = C_{\text{BANK } 1/2/3}$ $C = \text{SP} + A$ $\text{SP} = \text{SP} + A$ $C = \text{FP} + A$ $\text{FP} = \text{FP} + A$ </div>	$\langle \text{MEM_ACCESS_1} \rangle;$
Example:	<pre> if Z push r0, r5 = M[I0,M2]; if NZ pop DivResult; if Z SP = SP + r5; M[FP + rMAC] = B0; </pre>	
Type B:	<div> pushm $\langle \text{regList} \rangle$, $\text{SP} = \text{SP} + k_4$; $\text{SP} = \text{SP} - k_4$, popm $\langle \text{regList} \rangle$ $C_{\text{BANK } 1/2/3} = M[\text{FP} + k_{15}]$ $M[\text{FP} + k_{15}] = C_{\text{BANK } 1/2/3}$ $C_{\text{BANK } 1/2/3} = M[\text{SP} + k_{15}]$ $M[\text{SP} + k_{15}] = C_{\text{BANK } 1/2/3}$ $C = \text{SP} + k_{16}$ $\text{SP} = \text{SP} + k_{16}$ $C = \text{FP} + k_{16}$ $\text{FP} = \text{FP} + k_{16}$ push k_{16} push $C + k_{16}$ </div>	
Example:	<pre> push r0 + 4; pushm <r0, r7, rMACB>, SP = SP + 13; SP = SP - 11, popm <I0, I1, I2, I3, I4, I5, M0, L0>; I0 = M[SP + 0x123]; FP = FP + 0x500; </pre>	
Type C:	<div> push $C_{\text{BANK } 1/2/3}$ pop $C_{\text{BANK } 1/2/3}$ </div>	$\langle \text{MEM_ACCESS_1} \rangle \quad \langle \text{MEM_ACCESS_2} \rangle;$
Example:	<pre> push DivRemainder, r5 = M[I0,1], r1 = M[I4,-1]; </pre>	

Description:

Test the optional condition and, if TRUE, perform the specified stack operation. If the condition is FALSE, perform a NOP, but MEM_ACCESS_1 is still carried out. Omitting the condition performs the operation unconditionally. When using PUSH or POP, any single register from Bank 1, 2, or 3 can be used. When using PUSH Multiple or POP Multiple, all registers must be in the same bank (that is, all from Bank 1, 2, or 3).

NOTE `PUSH Multiple` and `POP Multiple` take one cycle per register pushed or popped.

Memory-mapped registers set the stack start address, frame pointer, and end address. The memory-mapped stack pointer is incremented with a `push` and decremented with a `pop`, and holds the current stack write address. A stack overflow or underflow causes a `STACK_FAULT` interrupt to occur, with the `STACK_OVERFLOW_PC` memory-mapped register being set to the PC where the most recent stack overflow or underflow event occurred.

Flags Generated:

All `PUSH/POP Multiple` Instructions:

All flags left unchanged.

`SP/FP Adjust` Instructions (for example, `SP = SP + 5`):

All flags set as per an `ADD` instruction (see Section 6.1).

`PUSH` Instructions:

Flags set depending on the value being pushed.

NOTE `V` and `C` can only be set with instructions such as `push r0 + 0x123`.

Z	Set if the value equals 0 and cleared otherwise	N	Set if the value is negative and cleared otherwise
V	Set if an arithmetic overflow occurs and cleared otherwise	C	Set if a carry is generated and cleared otherwise

`POP` Instructions:

Flags set depending on the value being popped.

Z	Set if the value equals 0 and cleared otherwise	N	Set if the value is negative and cleared otherwise
V	Cleared	C	Cleared

NOTE A Conditional `pop` with condition `FALSE` carries out the read regardless, but destination register, flags, and the memory-mapped register `STACK_POINTER` remain unchanged.

`push` of `DivResult` or `DivRemainder` while a divide is busy stalls program flow until the divide completes and then pushes the divide result. A conditional `push` of `DivResult` or `DivRemainder` with condition `FALSE` causes the program flow to stall anyway until the divide completes, but not push the result, set flags, or increment the stack pointer.

`pop` of `DivResult` or `DivRemainder` while the divider is busy fails to alter the value of `DivResult` or `DivRemainder`, but the stack pointer is decremented and flags are set accordingly.

When popping `DoLoopStart`, `DoLoopEnd`, or `r10`, all values must be popped at least two instructions before the end of a `do...loop`.

Writing to memory-mapped `M[$STACK_START_ADDR]` sets `M[$STACK_POINTER]` equal to `M[$STACK_START_ADDR]`. Popping to `rMAC (B) 12` sign-extends into `rMAC (B) 2`.

6.16 Program Flow: CALL, JUMP, RTS, RTI, and DO...LOOP

Syntax:

Type A:	<code><if cond></code>	<pre> jump A call A rts rti </pre>	<code>;</code>
Example:		<code>if Z jump r1;</code>	
Type B:	<code><if cond></code>	<pre> jump k₁₆ call k₁₆ do k₁₆ </pre>	<code>;</code>
Example:		<pre> r10 = 100; do loop; rMAC = rMAC + r0 * r1 r3 = M[I4,1]; loop: </pre>	

Description:

Omitting the condition performs the program flow. If the condition is `TRUE` perform the program flow described below. Otherwise, execute a `NOP`.

<code>jump</code>	Program execution jumps to the address in the operand, either a register or a constant (for example, an address label).
<code>do</code>	Starts the loop code execution. This instruction implements zero overhead loops called <i>do loop</i> . See Section 4.12.
<code>call</code>	Loads <code>rLink</code> register with return address (<code>PC+1</code>) and jumps to address in operand, either a register or a constant (for example, an address label).
<code>rti</code>	Sets <code>PC</code> equal to value of the memory-mapped <code>MM_RINTLINK</code> register, and restores flags to their pre-interrupt status (that is, bits 8 to 17 of <code>rFlags</code> register copied to bits 0 to 7).
<code>rts</code>	Sets <code>PC</code> equal to value of <code>rLink</code> register. Implement stack depths greater than 1 in software.

Flags Generated:

Flags `Z`, `N`, `V`, and `C` left unchanged.

NOTE Branches using an immediate `k16` value (no prefix) are relative. That is, the `k16` value (in the instruction word) is a signed offset from the current `PC` of where to branch to. The assembler (Kalasm3) will automatically calculate the relative value, and so the assembly syntax is always an immediate absolute address or label. Branches using a prefix (24-bits) and branches to a register location, e.g. `jump r0`, are interpreted as absolute branches by the processor.

6.17 Type A Miscellaneous One and Two Operand Instructions

Syntax:

Type A:	<if cond>	$C = A + 1$ $C = A - 1$ $C = \text{ABS } A$ $C = \text{MIN } A$ $C = \text{MAX } A$ $C = \text{ONEBITCOUNT } A$ $C = \text{TWOBITCOUNT } A$ $C = \text{MOD24 } A$ $C = A + 2$ $C = A + 4$ $C = A - 2$ $C = A - 4$ $C = \text{SE8 } A$ $C = \text{SE16 } A$	<MEM_ACCESS_1>;
----------------	------------------------	---	------------------------------

Example:

```
if Z r0 = MOD24 r1,
    r5 = M[I0,M2];
if NZ rMAC = rMACB - 1;
```

Description:

Test the optional condition and, if **TRUE**, perform the specified operation. If the condition is **FALSE**, perform a **NOP**, but **MEM_ACCESS_1** is still carried out. Omitting the condition performs the operation unconditionally. The operands can be any one of the 16 Bank 1 registers.

INC ($C = A + 1$) and **DEC** ($C = A - 1$) provide conditional increment and decrement instructions.

ABS provides a conditional absolute instruction, which multiplies negative numbers by -1.

MIN/MAX selects the signed smallest or largest operand from Reg C and Reg A, and places the result in Reg C.

ONEBITCOUNT accumulates together Reg A, assuming it represents 24 1-bit values. The result goes to Reg C. For example, if Reg A = 0x123456, after the instruction, Reg C = 9.

TWOBITCOUNT accumulates together Reg A, assuming it represents 12 2-bit values. The result goes to Reg C. For example, if Reg A = 0x123456, after the instruction, RegC = $0+1+0+2+0+3+1+0+1+1+1+2 = 12$.

MOD24 provides a conditional modulo 24 instruction. Reg A is interpreted as an unsigned number.

INC ($C = A + 2$), **INC** ($C = A + 4$), **INC** ($C = A - 2$), and **INC** ($C = A - 4$) provide conditional increment and decrement instructions

SE8 A and **SE16 A** provide a conditional instruction that takes the least significant 8 bits or 16 bits of Reg A and sign extends to Reg C.

Flags Generated:

For `INC`, `DEC`, and `ABS` instructions:

Z	Set if the value equals 0 and cleared otherwise	N	Set if the value is negative and cleared otherwise
V	Set if an arithmetic overflow occurs and cleared otherwise	C	Set if a carry is generated and cleared otherwise

For all other instructions:

Z	Set if the value equals 0 and cleared otherwise	N	Set if the value is negative and cleared otherwise
V	Cleared	C	Cleared

6.18 Indexed MEM_ACCESS_1 and MEM_ACCESS_2

Syntax:

	MEM_ACCESS_1		MEM_ACCESS_2
Type A:	Reg _{AG1} = M[I _{AG1} , M _{AG1}] M[I _{AG1} , M _{AG1}] = Reg _{AG1}	;	
Example:	M[I0, M0] = r1;		
Type C_{REG}:	Reg _{AG1} = M[I _{AG1} , M _{AG1}] M[I _{AG1} , M _{AG1}] = Reg _{AG1}		Reg _{AG2} = M[I _{AG2} , M _{AG2}] M[I _{AG2} , M _{AG2}] = Reg _{AG2}
Type C_{CONST}:	Reg _{AG1} = M[I _{AG1} , MK _{AG1}] M[I _{AG1} , MK _{AG1}] = Reg _{AG1}		Reg _{AG2} = M[I _{AG2} , MK _{AG2}] M[I _{AG2} , MK _{AG2}] = Reg _{AG2}
Example:	r0 = M[I0, M0], M[I4, M1] = r1;		

Permitted Registers

Reg _{AG1} /Reg _{AG2}	r0, r1, r2, r3, r4, r5, and rMAC	I _{AG1}	I0, I1, I2, and I3
M _{AG1} /M _{AG2}	M0, M1, M2, and M3	I _{AG2}	I4, I5, I6, and I7
MK _{AG1} /MK _{AG2}	-1, 0, 1 and 2		

Description:

Any Type C instruction also can perform up to two memory reads or writes in the same instruction cycle as the main ALU part of the instruction. Reg_{AG1/AG2} selects the source or destination register for the memory read or write. I_{AG1/AG2} selects the index register to use for the memory read. M_{AG1/AG2} selects the modify register, or MK_{AG1/AG2} selects the modify constant (-1, 0, +1, or +2) to use for the post modify of the index register.

Flags Generated:

No flags are affected by the memory access part of instructions.

NOTE Reg_{AG1/AG2} selects one of the first eight Bank 1 registers (that is Null, rMAC, or r0 to r5). If the Null register is selected, then no memory read or write is performed.

Type A instructions use AG1, so only index registers I0 to I3 can be used. They must use a modify register rather than a modify constant.

Type C instructions use AG1 and AG2 so one memory access must use one of I0 to I3 and the other must use one of I4 to I7. They must either both use a modify register or both use a modify constant.

The index register post-modify takes place after the instruction has executed. The modify register or constant does not affect the address that is used in the memory access or the result of any calculation in the main instruction involving the index register.

Modulo indexing, controlled by the length and base registers, is applied during the post-modify part of an indexed memory access.

7 Instruction Coding

Instruction coding is relatively simple and orthogonal so that the instruction decode is efficient. There are three basic coding formats, type A, type B, and type C. See Section 7.1 to 7.20.

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	Type		
OP_CODE							0	0	RegC MultAddSubSelect				RegA						AG1 Write	RegAG1			IAG1		MAG1		RegB			cond			A	
							0	1											K16										B					
							1	0																					MAG2		CREG			
							1	1																					MKAG2		CCONST			
OP_CODE (MAC with ADD/SUB)							1	ADD/SUB	RegD	RegC'		RegA						AG1 Write	RegAG1			IAG1		MAG1		RegAG2			IAG2		MAG2		CMAC_ADDSUB	
1111010							1	RegC										MAG1		AG2 Write			RegAG2			IAG2		MKAG2		CCONST				
							MAG1											RegB			cond			ASubWord										
							K11											BSubWord																
111111010000																	KPREFIX[20:0]										PFIX							

Figure 7-1 Instruction Coding Format

Type	Instruction type (see Section 7.1 through 7.3)
OP_CODE	Selects the instruction operation. (See Section 7.5.)
RegA _{B1/2}	Selects a register for the first source operand for instructions. Bank 1 registers are the default. Bank 2 is selected in OP_CODE for certain instructions, indicated by _{B1/2} .
RegB _{B1/2}	Selects a register for the second source operand for instructions. Bank 1 registers are the default. Bank 2 is selected in OP_CODE for certain instructions, indicated by _{B1/2} .
RegC _{B1/2/3}	Selects 1 of 16 registers for the destination register for instructions. For type C instructions, RegC also defines the first source operand (see Table 7-1). Bank 1 registers are used by default. Bank 2 or Bank 3 is selected in OP_CODE or StackBankSelect for certain instructions, indicated by _{B1/2/3} .
cond	Selects an optional condition to be met for the instruction to be executed. Otherwise, a NOP instruction is executed (see Section 4.5.10).
k ₁₆ /k _{PREFIX}	A 16-bit or an 8-bit constant used by type B instructions.
AG1 / AG2	Selects whether the indexed memory access is a read (0) or a write (1).
RegAG1/RegAG2	Selects one of the first eight Bank 1 registers (null, rMAC, r0 to r5) for the source or destination register of multifunction memory reads and writes. If Null is selected, then no read or write is performed.
I _{AG1} /I _{AG2}	Selects one of the index registers, I0 to I3 for AG1 and I4 to I7 for AG2, for multifunction memory reads and writes.
M _{AG1} /M _{AG2}	Selects modify register M0 to M3 for multifunction memory reads and writes.
MK _{AG1} /MK _{AG2}	Selects a fixed constant to use for the modify: -1, 0, +1, or +2
StackBankSelect	Selects the bank of registers to use for stack operations.

NOTE Throughout, *OP* refers to *operation*, which can include addition (+), subtraction (-), multiplication (*), AND, OR, and exclusive OR (XOR).

7.1 Type A Instruction

Type A is a conditional instruction with an additional single memory read or write operation. The instruction can accept two input operands and one output operand (which can be different).

Operands can be any of the 16 Bank 1 registers. (LOAD/STORE instructions permit the use of Bank 2 registers. PUSH/POP instructions permit the use of Bank 1, Bank 2, and special Bank 3 registers.)

For the memory access, index registers *I0* to *I3* select the address. The destination or source register can be any of the first eight Bank 1 registers except *Null*: *rMAC* or *r0* to *r5*. At the end of the instruction, the index register is post-modified by one of the modify registers *M0* to *M3*.

A type A instruction has the following syntax:

```
<if cond> RegC = RegA OP RegB <MEM_ACCESS_1>;
```

NOTE Program flow instructions, such as *jump*, *call*, and *rts*, use a slight variation on types A, B, and C, because they can always be conditional.

7.2 Type B Instruction

Type B is a non-conditional instruction similar to type A, but with one of the operands being a 16-bit constant stored in the instruction word.

To use a 24-bit constant, prefix the type B instruction with the prefix (PFI) instruction. No additional memory access operation is permitted. The PFI instruction is added automatically by the assembler *kalasm3*, if needed.

A type B instruction has the following syntax:

```
RegC = RegA OP constant;
```

NOTE Program flow instructions, such as *jump*, *call*, and *rts*, use a slight variation on types A, B, and C, because they can always be conditional.

7.3 Type C Instruction

Type C is a non-conditional instruction similar to a type A, except that one of the input operands also is the output operand.

In addition, two memory accesses (reads, writes, or a combination) can occur in the same clock cycle. One memory access uses *AG1* (index registers *I0* to *I3*). The other uses *AG2* (index registers *I4* to *I7*).

The index registers can be post-modified by the *M0* to *M3* registers (this is a type *C_{REG}* instruction) or by a 2-bit signed modify constant where the valid values are -1, 0, 1, or 2 (this is a type *C_{CONST}* instruction).

A type C instruction has the following syntax:

```
RegC = RegC OP RegA <MEM_ACCESS_1> <MEM_ACCESS_2>;
```

NOTE Program flow instructions, such as `jump`, `call`, and `rts`, use a slight variation on types A, B, and C, because they can always be conditional.

7.4 Special Cases

Program flow instructions, such as `jump`, `call`, `rts`, etc. use a slight variation on the above types, because they can always be conditional.

7.5 OP_CODE Coding

The `OP_CODE` field is decoded according to the instruction code type (A, B, or C).

[Table 7-1](#) describes the instruction decoding of the `OP_CODE` field in Figure 7-1.

Table 7-1 OPCODE Coding Format

OP_CODE				Action (Type A) ^a	Action (Type B)	Action (Type C/REG/CONST) ^b	Description
000	AM		C	$\text{RegC} = \text{RegA} + \text{RegB}$	$\text{RegC} = \text{RegA} + k_{16}$	$\text{RegC} = \text{RegC} + \text{RegA}$	Add ^c
001	AM		C	$\text{RegC} = \text{RegA} - \text{RegB}$	$\text{RegC} = \text{RegA} - k_{16}$	$\text{RegC} = \text{RegC} - \text{RegA}$	Subtract
010	B2RS			$\text{RegC}_{B1/2} = \text{RegA}_{B1/2} + \text{RegB}_{B1/2}$	$\text{RegC}_{B1/2} = \text{RegA}_{B1/2} + k_{16}$	$\text{RegC}_{B1/2} = \text{RegC}_{B1/2} + \text{RegA}_{B1/2}$	Bank 1/2 Add
011	B2RS			$\text{RegC}_{B1/2} = \text{RegA}_{B1/2} - \text{RegB}_{B1/2}$	$\text{RegC}_{B1/2} = \text{RegA}_{B1/2} - k_{16}$ $\text{RegC}_{B1/2} = k_{16} - \text{RegA}_{B1/2}$	$\text{RegC}_{B1/2} = \text{RegC}_{B1/2} - \text{RegA}_{B1/2}$	Bank 1/2 Subtract
100	0	0	0	$\text{RegC} = \text{RegA} \text{ AND } \text{RegB}$	$\text{RegC} = \text{RegA} \text{ AND } k_{16}$	$\text{RegC} = \text{RegC} \text{ AND } \text{RegA}$	Logical AND
100	0	0	1	$\text{RegC} = \text{RegA} \text{ OR } \text{RegB}$	$\text{RegC} = \text{RegA} \text{ OR } k_{16}$	$\text{RegC} = \text{RegC} \text{ OR } \text{RegA}$	Logical OR
100	0	1	0	$\text{RegC} = \text{RegA} \text{ XOR } \text{RegB}$	$\text{RegC} = \text{RegA} \text{ XOR } k_{16}$	$\text{RegC} = \text{RegC} \text{ XOR } \text{RegA}$	Logical XOR
100	0	1	1	$\text{RegC} = \text{RegA} \text{ LSHIFT } \text{RegB}$	$\text{RegC} = \text{RegA} \text{ LSHIFT } k_{16}$	$\text{RegC} = \text{RegC} \text{ LSHIFT } \text{RegA}$	Logical Shift
100	1	0	0	$\text{RegC} = \text{RegA} \text{ ASHIFT } \text{RegB}$	$\text{RegC} = \text{RegA} \text{ ASHIFT } k_{16}$	$\text{RegC} = \text{RegC} \text{ ASHIFT } \text{RegA}$	Arithmetic Shift
100	1	1	V	$\text{RegC} = \text{RegA} * \text{RegB}$ (int)	$\text{RegC} = \text{RegA} * k_{16}$ (int)	$\text{RegC} = \text{RegC} * \text{RegA}$ (int)	Integer signed multiply
100	1	0	1	$\text{RegC} = \text{RegA} * \text{RegB}$ (frac)	$\text{RegC} = \text{RegA} * k_{16}$ (frac)	$\text{RegC} = \text{RegC} * \text{RegA}$ (frac)	Fractional signed multiply
101	0	S	S	$\text{rMAC}(\text{B}) = \text{rMAC}(\text{B}) + \text{RegA} * \text{RegB}$ [$\text{r0} = \text{RegD} \pm \text{rMAC}(/B)$]	$\text{rMAC}(\text{B}) = \text{rMAC}(\text{B}) + \text{RegA} * k_{16}$ [$\text{r0} = \text{RegD} \pm \text{rMAC}(/B)$]	$\text{rMAC} = \text{rMAC} + \text{RegC} * \text{RegA}$	Multiply accumulate (56-bit) with optional ADD/SUB

Table 7-1 OPCODE Coding Format (cont.)

OP_CODE				Action (Type A) ^a	Action (Type B)	Action (Type CREG/CONST) ^b	Description
101	1	S	S	$rMAC(B) = rMAC(B) - RegA * RegB$ $[r0 = RegD \pm rMAC(/B)]$	$rMAC(B) = rMAC(B) - RegA * k_{16} [r0 = RegD \pm rMAC(/B)]$	$rMAC = rMAC - RegC * RegA$	Multiply subtract (56-bit) with optional ADD/SUB
110	0	S	S	$rMAC(B) = RegA * RegB$ $[r0 = RegD \pm rMAC(/B)]$	$rMAC(B) = RegA * k_{16} [r0 = RegD \pm rMAC(/B)]$	$rMAC = RegC * RegA$	Multiply (48-bit) with optional ADD/SUB
110	1	0	0	$RegC = M[RegA + RegB]$	$RegC = M[RegA + k_{16}]$	$RegC = M[RegC + RegA]$	Load with offset
110	1	0	1	$M[RegA + RegB] = RegC$	$M[RegA + k_{16}] = RegC$	$rMAC = rMAC + RegC' * RegA$ $r0 = RegD \pm rMACB$	Store with offset, mult-acc (56-bit) with ADD/SUB
110	1	1	0	$RegC = SignDet RegA$	$Div = rMAC/RegA$ $RegC = DivResult$ $RegC = DivRemainder$	$RegC = BlkSignDet RegA$	Sign detect, divide, block sign detect
110	1	1	1	JUMP RegA (RTS if $RegA=rLink$) (RTI if $RegA=rFlags$)	if $[RegC=cond]$ JUMP k_{16}	$rMAC = rMAC - RegC' * RegA$ $r0 = RegD \pm rMACB$	Jump to program address, return from subroutine, return from interrupt, multi-sub (56-bit) with ADD/SUB
111	0	0	0	CALL RegA	if $[RegC=cond]$ CALL k_{16}	$rMAC = RegC' * RegA$ $r0 = RegD \pm rMACB$	Call subroutine, return from interrupt, multiply (56-bit) with ADD/SUB
111	0	0	1	$RegC = MIN/MAX/ABS...RegA$	DO ... LOOP	$rMACB = rMACB + RegC' * RegA$ $r0 = RegD \pm rMAC$	Misc 1 and 2 Operand instructions, Do Loop, multi-acc (56-bit) with ADD/SUB
111	0	1	0	Future use	$RegC = k_{16} LSHIFT RegA$	$rMACB = rMACB - RegC' * RegA$ $r0 = RegD \pm rMAC$	Logical Shift of a constant, multi-sub (56-bit) with ADD/SUB
111	0	1	1	Future use	$RegC = k_{16} ASHIFT RegA$	$rMACB = RegC' * RegA$ $r0 = RegD \pm rMAC$	Arithmetic Shift of a constant, multiply (56-bit) with ADD/SUB

Table 7-1 OPCODE Coding Format (cont.)

OP_CODE				Action (Type A) ^a	Action (Type B)	Action (Type C _{REG/CONST}) ^b	Description
111	1	0	0	PUSH RegC _{B1/2/3} POP RegC _{B1/2/3} LOAD/STORE to stack with offset Stack/frame pointer adjust	PUSH Multiple POP Multiple LOAD/STORE to stack with offset Stack/frame pointer adjust	PUSH RegC _{B1/2/3} POP RegC _{B1/2/3}	New stack operations
111	1	0	1	RegC _{B1/2} = Mxx[RegA _{B1/2} ± RegB] Myy[RegA _{B1/2} ± RegB] = RegC _{B1/2}	RegC _{B1/2} = Mxx[RegA _{B1/2} ± k ₁₁] Myy[RegA _{B1/2} ± k ₁₁] = RegC _{B1/2}	Not available (used by type- A/B)	Subword memory read/ write instructions
111	1	1	0	Future use			
111	1	1	1		PREFIX instruction		Enables 24-bit constants for Type B operations—by prefixing the following instruction's k ₁₆ by the 8-bit, k _{PREFIX} value.

^a Type A is conditional with if[cond] and can contain single memory access MEM_ACCESS_1.

^b Type C permits 2 simultaneous memory accesses MEM_ACCESS_1 and MEM_ACCESS_2.

^c The add instruction is also used to implement LOAD/STORE to registers or memory by setting one of the source registers as Null. Setting all three operands as Null implements a no-operation (NOP) instruction.

See Section 7.11 and Section 7.16 for encoding of k₁₆ for rMAC shift instructions and DIVIDE instructions.

7.6 AM Field

The AM Field selects different memory addressing modes.

Table 7-2 AM Field

AM	Type A	Type B	Type C _{REG/CONST}
00	RegC = RegA [OP] RegB	RegC = RegA [OP] k ₁₆	RegC = RegC [OP] RegA
01	RegC = RegA [OP] M[RegB]	RegC = RegA [OP] M[k ₁₆]	RegC = RegC [OP] M[RegA]
10	RegC = M[RegA] [OP] RegB	RegC = M[RegA] [OP] k ₁₆	–
11	M[RegC] = RegA [OP] RegB	M[k ₁₆] = RegC [OP] RegA	–

7.7 Carry Field (C Field)

This field selects whether addition or subtraction is performed with carry and the appropriate state.

Table 7-3 Field Options

c	Description
0	Do not use carry or borrow
1	Use carry or borrow

7.8 Bank 1/2 Register Select Field (B2RS Field)

Subtract from constant type B instructions are valid only for SUBTRACT opcodes.

Table 7-4 B2RS Field

B2RS	Type A	Type B	Type C _{REG/CONST}
000	$\text{RegC}_{\text{Bank1}} = \text{RegA}_{\text{Bank1}} \pm \text{RegB}_{\text{Bank1}}$	$\text{RegC}_{\text{Bank1}} = \text{RegA}_{\text{Bank1}} \pm K_{16}$	$\text{RegC}_{\text{Bank1}} = \text{RegC}_{\text{Bank1}} \pm \text{RegA}_{\text{Bank1}}$
001	$\text{RegC}_{\text{Bank1}} = \text{RegA}_{\text{Bank1}} \pm \text{RegB}_{\text{Bank2}}$	$\text{RegC}_{\text{Bank1}} = K_{16} - \text{RegA}_{\text{Bank1}}$	$\text{RegC}_{\text{Bank1}} = \text{RegC}_{\text{Bank1}} \pm \text{RegA}_{\text{Bank2}}$
010	$\text{RegC}_{\text{Bank1}} = \text{RegA}_{\text{Bank2}} \pm \text{RegB}_{\text{Bank1}}$	$\text{RegC}_{\text{Bank1}} = \text{RegA}_{\text{Bank2}} \pm K_{16}$	
011	$\text{RegC}_{\text{Bank1}} = \text{RegA}_{\text{Bank2}} \pm \text{RegB}_{\text{Bank2}}$	$\text{RegC}_{\text{Bank1}} = K_{16} - \text{RegA}_{\text{Bank2}}$	
100	$\text{RegC}_{\text{Bank2}} = \text{RegA}_{\text{Bank1}} \pm \text{RegB}_{\text{Bank1}}$	$\text{RegC}_{\text{Bank2}} = \text{RegA}_{\text{Bank1}} \pm K_{16}$	
101	$\text{RegC}_{\text{Bank2}} = \text{RegA}_{\text{Bank1}} \pm \text{RegB}_{\text{Bank2}}$	$\text{RegC}_{\text{Bank2}} = K_{16} - \text{RegA}_{\text{Bank1}}$	
110	$\text{RegC}_{\text{Bank2}} = \text{RegA}_{\text{Bank2}} \pm \text{RegB}_{\text{Bank1}}$	$\text{RegC}_{\text{Bank2}} = \text{RegA}_{\text{Bank2}} \pm K_{16}$	$\text{RegC}_{\text{Bank2}} = \text{RegC}_{\text{Bank2}} \pm \text{RegA}_{\text{Bank1}}$
111	$\text{RegC}_{\text{Bank2}} = \text{RegA}_{\text{Bank2}} \pm \text{RegB}_{\text{Bank2}}$	$\text{RegC}_{\text{Bank2}} = K_{16} - \text{RegA}_{\text{Bank2}}$	$\text{RegC}_{\text{Bank2}} = \text{RegC}_{\text{Bank2}} \pm \text{RegA}_{\text{Bank2}}$

7.9 Saturation Select Field (V Field)

This field selects whether to enable saturation on the result.

NOTE The addition and subtraction instructions also can saturate on overflow by setting the ARITHMETIC_MODE memory mapped register bit.

Table 7-5 V Field

v	Description
0	No saturation
1	Saturation

7.10 Sign Select Field (S Field)

This field selects signed and unsigned multiples.

Table 7-6 S Field

s	Description
00	Unsigned x unsigned
01	Unsigned x signed
10	Signed x unsigned
11	Signed x signed

7.11 k₁₆ Coding for LSHIFT and ASHIFT

The k₁₆ coding section from the instruction coding format splits into its individual bits (see Figure 7-1).

[Table 7-7](#) describes their functionality.

Table 7-7 k₁₆ Coding Shift Format

Bit															
15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	0	0	0	Dest_Sel			ShiftAmount						

ShiftAmount is the amount to shift the input, it is a signed 7-bit number.

Positive is a left shift. Negative is a right shift.

Dest_Sel is when the destination register is rMAC and selects how the 24-bit output from the shifter is used.

7.12 rMAC Sub-registers

The full 56-bits of the rMAC register are accessible as individual sub-registers. See Section 7.13 and Section 7.14 for information about how to load the individual subregisters.

Table 7-8 rMAC Sub-registers

Bits	[55:48]	[47:24]	[23:0]
rMAC Register	rMAC2	rMAC1	rMAC0

7.13 ASHIFT

The arithmetic shift instruction is encoded within the instruction-coding format (see Figure 7-1). See [LSHIFT and ASHIFT \(56 bit\)](#) for information about the ASHIFT instruction.

NOTE The LSHIFT and ASHIFT instructions enable the user to write to the individual subregisters (rMAC2, rMAC1, and rMAC0), providing a way of loading the rMAC with a double precision number. Shift operations of the rMAC also are enabled, which speeds up double-precision calculations. The destination is the rMAC register.

Table 7-9 ASHIFT

Dest_Sel	New rMAC2	New rMAC1	New rMAC0	Example
001	Sign extend	Sign extend	SHIFTER_OUTPUT	rMAC = r? ASHIFT k ₁₆ (LO)
000 (rFlags)	Sign extend	SHIFTER_OUTPUT	Trailing zeros	rMAC = r? ASHIFT k ₁₆ (MI)
000 (rMAC (B))	SHIFTER_OUTPUT			rMAC (B) = r? ASHIFT k ₁₆ (56 bit)
010	SHIFTER_OUTPUT	Trailing zeros	Trailing zeros	rMAC = r? ASHIFT k ₁₆ (HI)
101	Old rMAC2	Old rMAC1	SHIFTER_OUTPUT	rMAC0 = r? ASHIFT k ₁₆
100	Sign extend	SHIFTER_OUTPUT	Old rMAC0	rMAC12 = r? ASHIFT k ₁₆
110	SHIFTER_OUTPUT	Old rMAC1	Old rMAC0	rMAC2 = r? ASHIFT k ₁₆

7.14 LSHIFT

The logical shift instruction is encoded within the instruction-coding format (see Figure 7-1). See [LSHIFT and ASHIFT \(56 bit\)](#) for information about the LSHIFT instruction.

NOTE The LSHIFT and ASHIFT instructions enable the user to write to the individual subregisters (rMAC2, rMAC1, and rMAC0), providing a way of loading the rMAC with a double precision number. Shift operations of the rMAC also are enabled, which speeds up double-precision calculations. The destination is the rMAC register.

Table 7-10 LSHIFT

Dest_Sel (RegC)	New rMAC2	New rMAC1	New rMAC0	Example
001	Zero fill	Zero fill	SHIFTER_OUTPUT	rMAC (B) = r? LSHIFT k ₁₆ (LO)
000 (rFlags)	Zero fill	SHIFTER_OUTPUT	Trailing zeros	rMAC (B) = r? LSHIFT k ₁₆ (MI)
000 (rMAC (B))	SHIFTER_OUTPUT			rMAC (B) = r? LSHIFT k ₁₆ (56 bit)

Table 7-10 LSHIFT (cont.)

Dest_Sel (RegC)	New rMAC2	New rMAC1	New rMAC0	Example
010	SHIFTER_OUTPUT	Trailing zeros	Trailing zeros	$rMAC(B) = r? \text{ LSHIFT } k_{16} \text{ (HI)}$
101	Old rMAC2	Old rMAC1	SHIFTER_OUTPUT	$rMAC(B)0 = r? \text{ LSHIFT } k_{16}$
100	Zero fill	SHIFTER_OUTPUT	Old rMAC0	$rMAC(B)12 = r? \text{ LSHIFT } k_{16}$
110	SHIFTER_OUTPUT	Old rMAC1	Old rMAC0	$rMAC(B)2 = r? \text{ LSHIFT } k_{16}$

7.15 Type A and Type C LSHIFT56 and ASHIFT56

RegC	Shift Destination Register	Data Width of Shift Output (bit)
rFlags	rMAC	24
rMAC	rMAC	56
rMACB	rMACB	56

7.16 k_{16} Coding Divide Instructions

Table 7-11 describes the k_{16} coding divide instruction and shows how the background divide instruction is encoded within the instruction-coding format (see Figure7-1).

Table 7-11 Divide Field

Bit															
15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Not Used														Div	

Table 7-12 describes the functionality of the individual divide bits.

Table 7-12 Divide Field States

DIV	Assembly Syntax	Operation
00	Div = rMAC/RegA	Initiates a 48-bit/24-bit divide using rMAC as the dividend and RegA as the divisor.
01	RegC = DivResult	There must be at least 12 cycles between the divide instruction and using the result and/or remainder. In this period, normal program execution continues. If the result request is early, then program execution automatically delays until the result is available.
10	RegC = DivRemainder	
11	Div = rMACB/RegA	Initiates a 48-bit/24-bit divide using rMACB as the dividend and RegA as the divisor.

7.17 Type A Miscellaneous One and Two Operand Instruction Encodings

Table 7-13 Encodings for One and Two Operand Instructions

RegA	RegC	RegB	Operation
RegA	RegC	0 0 1 0	RegC = RegA + 1
RegA	RegC	0 0 1 1	RegC = RegA - 1
RegA	RegC	0 1 0 0	RegC = ABS RegA
RegA	RegC	0 1 0 1	RegC = MIN RegA
RegA	RegC	0 1 1 0	RegC = MAX RegA
RegA	RegC	0 1 1 1	RegC = TWOBITCOUNT RegA
RegA	RegC	1 0 0 0	RegC = MOD24 RegA
RegA	RegC	1 0 0 1	RegC = ONEBITCOUNT RegA
RegA	RegC	1 0 1 0	RegC = RegA + 2
RegA	RegC	1 0 1 1	RegC = RegA + 4
RegA	RegC	1 1 0 0	RegC = RegA - 2
RegA	RegC	1 1 0 1	RegC = RegA - 4
RegA	RegC	1 1 1 0	RegC = SE8 RegA
RegA	RegC	1 1 1 1	RegC = SE16 RegA

7.18 Type A Stack Encoding

Table 7-14 Type A Stack Operation Encoding

Instruction	RegC	RegA				RegB
push rStack;	RegSelect	0	0	BankSelect		Unused (set to 0)
pop rStack;		0	1			
rStack = M[FP + RegOffset];		1	0			
M[FP + RegOffset] = rStack;		1	1			
RegC = SP + RegOffset; OR, if (RegC==Null) SP = SP + RegOffset;	RegC	0	0	1	1	RegOffset
RegC = FP + RegOffset; OR, if (RegC==Null) FP = FP + RegOffset;		0	1	1	1	

NOTE rStack is the register selected by RegSelect in the bank selected by BankSelect.

- 0: Bank 1
- 1: Bank 2
- 2: Bank 3

FP is the frame pointer.

SP is the stack pointer.

7.19 Type B Stack Encoding

Table 7-15 Type B Stack Operation Encoding

Instruction	RegC	RegA				K ₁₆
pushm <regList>, SP = SP + StackAdjust;	StackAdjust (unsigned)	0	0	BankSelect		RegSelectBitfield
SP = SP - StackAdjust, popm <reglist>;	StackAdjust (unsigned)	0	1			
rStack = M[FP + Offset];	RegSelect	1	0			0
rStack = M[SP + Offset];						1
M[FP + Offset] = rStack;		1	1			0
M[SP + Offset] = rStack;						1
RegC = SP + Offset; OR, if (RegC==Null) SP = SP + Offset;	RegC	0	0	1	1	Offset (signed)
RegC = FP + Offset; OR, if (RegC==Null) FP = FP + Offset;		0	1	1	1	
push RegC + k16	RegC (Bank1 only)	1	1	1	1	Value (signed)

NOTE rStack is the register selected by RegSelect in the bank selected by BankSelect.

- 0: Bank 1
- 1: Bank 2
- 2: Bank 3

FP is the frame pointer.

SP is the stack pointer.

StackAdjust is a 4-bit unsigned value of how much to adjust SP by after a pushm (increment) or before a popm (decrement).

RegSelectBitfield has a bit per register, if a bit is set then that register (in the bank selected by BankSelect) is pushed/popped. This provides a push/pop multiple encoding.

Table 7-16 BankSelect = 0 (Bank 1), or BankSelect = 1 (Bank 2)

Bit	BankSelect = 0 (Bank 1)	BankSelect = 1 (Bank 2)	BankSelect = 2 (Bank 3)
	Register	Register	Register
0	FP (=SP) (Special case: after pushing, FP is set equal to the value of SP before the pushm instruction)	I0	rMAC2
1	rMAC (24-bit)	I1	rMAC12
2	r0	I2	rMAC0
3	r1	I3	DoLoopStart
4	r2	I4	DoLoopEnd
5	r3	I5	DivResult
6	r4	I6	DivRemainder
7	r5	I7	rMACB2
8	r6	M0	rMACB12
9	r7	M1	rMACB0
10	r8	M2	B0
11	r9	M3	B1
12	r10	L0	B4
13	rLink	L1	B5
14	rFlags	L4	FP
15	rMACB	L5	SP

7.20 Type C Stack Encoding

Table 7-17 Type C Stack Operation Encoding

Instruction	RegC	RegA		
push rStack;	RegSelect	0	0	BankSelect
pop rStack;		0	1	

NOTE rStack is the register selected by RegSelect in the bank selected by BankSelect.

- 0: Bank 1
- 1: Bank 2
- 2: Bank 3

7.21 Type C MAC with ADD/SUB Instruction

Table 7-18 BlueCore7/BlueCore8 Type C MAC with ADD/SUB Encoding

Instruction	OP_CODE	25	24	23	22	21	20	19	18	17	16	15:0	Type
$rMAC = rMAC + RegC' * RegA, r0 = RegD + rMACB$	1 1 0 1 0 1	1	0	RegD	RegC'							2 indexed memory accesses (of type C_{REG} only, that is, with modify registers rather than constants)	C_{MAC_ADDSUB}
$rMAC = rMAC + RegC' * RegA, r0 = RegD - rMACB$			1										
$rMAC = rMAC - RegC' * RegA, r0 = RegD + rMACB$	1 1 0 1 1 1	1	0										
$rMAC = rMAC - RegC' * RegA, r0 = RegD - rMACB$			1										
$rMAC = RegC' * RegA, r0 = RegD + rMACB$	1 1 1 0 0 0	1	0										
$rMAC = RegC' * RegA, r0 = RegD - rMACB$			1										
$rMACB = rMACB + RegC' * RegA, r0 = RegD + rMAC$	1 1 1 0 0 1	1	0										
$rMACB = rMACB + RegC' * RegA, r0 = RegD - rMAC$			1										
$rMACB = rMACB - RegC' * RegA, r0 = RegD + rMAC$	1 1 1 0 1 0	1	0										
$rMACB = rMACB - RegC' * RegA, r0 = RegD - rMAC$			1										
$rMACB = RegC' * RegA, r0 = RegD + rMAC$	1 1 1 0 1 1	1	0										
$rMACB = RegC' * RegA, r0 = RegD - rMAC$			1										

NOTE RegC' specifies the first multiplicand. This is a 3-bit value that selects rMACB (special case when $\text{RegC}' = 0$), rMAC , r0 , r1 , r2 , r3 , r4 , or r5 .

RegA specifies the second multiplicand. This is a standard 4-bit Bank 1 register select value.

RegD specifies the first operand of the add/sub operation. It is a 1-bit value and selects either r1 (0) or r2 (1).

The second operand of the add/sub is always an rMAC (B) register and is the opposite rMAC (B) register to that used for the multiply part.

7.22 Type A and B MAC and Optional ADD/SUB Instruction

Table 7-19 BlueCore7/BlueCore8 Type A and B MAC and Optional ADD/SUB Encoding

Instruction	OP_CODE			RegC			RegA	RegB/k ₁₆
RegC' = RegC' + RegA * RegB;	1 0 1 0	S	S	0	RegC'		RegA 1 st multiplicand	RegB/k ₁₆ 2 nd multiplicand
RegC'' = RegC'' + RegA * RegB r0 = RegD + RegE		Reg D	0	1	Reg E	Reg C''		
RegC'' = RegC'' + RegA * RegB r0 = RegD - RegE			1	1				

NOTE RegC' specifies the accumulator register (for when there is no supplementary ADD/SUB operation). This is a 3-bit value that selects rMACB (special case when $\text{RegC}' = 0$), rMAC , r0 , r1 , r2 , r3 , r4 , or r5 .

RegC'' specifies the accumulator register (for when there is a supplementary ADD/SUB operation). This is a 2-bit value that selects rMACB (special case when $\text{RegC}' = 0$), rMAC , r1 , or r2 .

RegD specifies the first operand of the add/sub operation. It is a 1-bit value and selects either r1 (0) or r2 (1).

RegE specifies the second operand of the ADD/SUB operation. It is a 1-bit value and selects either rMAC (0) or rMACB (1).

8 Kalimba DSP Peripherals

The Kalimba DSP for BlueCore consists of the DSP core, the DSP peripherals, and their associated interfaces.

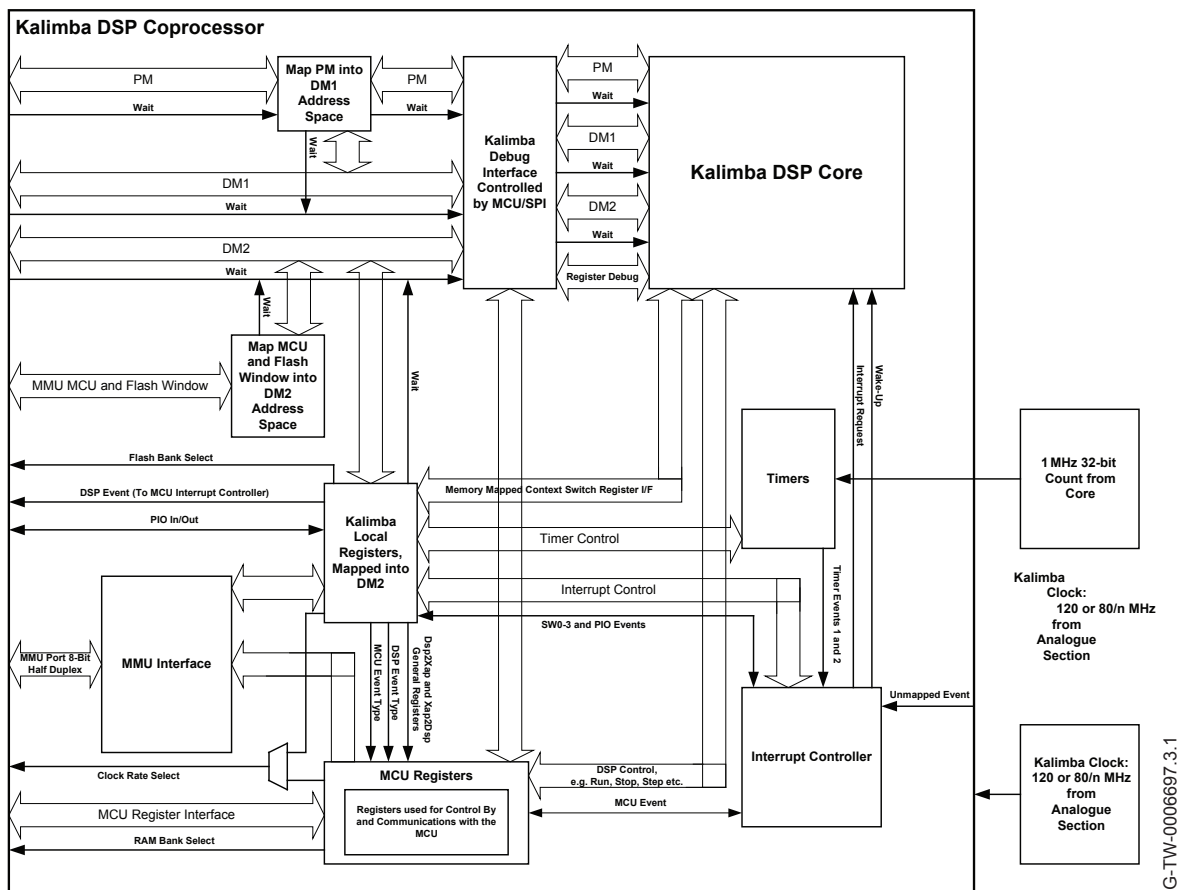


Figure 8-1 Kalimba DSP Peripheral Interfaces

The Kalimba DSP peripherals include:

- MMU interface, for stream transfers to and from the rest of the IC
- Three memory-mapped windows into the flash or ROM
- Two memory-mapped windows into the MCU RAM
- Memory-mapped register interface into the MCU RAM and I/O map
- Program memory window into flash or ROM with 1024-word or 512-word, direct-mapped cache

- Two 1 μ s-resolution timers
- Interrupt controller with three priority levels
- Memory-mapped access to the DSP program memory through DM1
- Clock rate divider controllable by either the DSP or the MCU
- Debug interface

8.1 MMU Interface

The Kalimba DSP MMU interface contains 12 virtual read ports and 12 virtual write ports.




Figure 8-2 Example of MMU Interface Usage for a Wireless MP3 Player

8.1.1 Read Ports

Twelve virtual read ports appear as memory-mapped registers in DM2.

The ports have the ability to use the following:

- 8-bit, 16-bit, or 24-bit word size
- Byte swap capability (little endian/big endian)
- Sign-extension, if required

8.1.2 Write Ports

Twelve virtual write ports appear as memory-mapped registers in DM2.

The ports have the ability to use the following:

- 8-bit, 16-bit, or 24-bit word size
- Byte swap capability (little endian/big endian)
- Sign-extension, if required

8.2 DSP Timers

The Kalimba DSP timer feature uses a 32-bit `TIMER_TIME` read-only register.

The Kalimba DSP timer features are:

- A 32-bit `TIMER_TIME` register (read-only) clocked at 1 MHz
- Two trigger value registers. Each trigger value register, if enabled, causes an appropriate interrupt (LS 24 bits of `TIMER_TIME` only used for trigger values).

G-TW-0006698.2.2

8.3 Kalimba Interrupt Controller

The interrupt controller has many interrupt lines as inputs. The interrupt controller monitors these lines looking for rising edges.

Monitoring of the lines is controlled by `INT_SOURCES_EN`. The only time that the lines are not monitored is if the interrupt controller is disabled (`INT_GBL_ENABLE == 0`). The lines are monitored even when interrupts are blocked (`INT_UNBLOCK == 0`).

When a rising edge is detected from an interrupt source, it is called a *registered interrupt*.

NOTE Multiple interrupt lines can be registered at any one time.

8.3.1 Interrupt Controller Functionality

The functionality of the interrupt controller includes interrupt sources, priority levels, registers to control the interrupt controller, and an event signal.

Interrupt controller functionality includes the following:

- Selectable interrupt sources:
 - ☐ Timer1 and Timer2
 - ☐ MCU event
 - ☐ PIO event
 - ☐ Buffer error or software error
 - ☐ `SwEvent0`, `SwEvent1`, `SwEvent2`, and `SwEvent3`
- Three interrupt priority levels
- Registers to save and restore the interrupt controller which enables nested interrupts
- Optional event signal to cause clock rate change

8.3.2 Registered Interrupt Processing

If the interrupt controller finds that it has a registered interrupt, it has two options: The DSP core is in the process of handling a previous interrupt and running the early part of the ISR code up until the point of writing to the `INT_ACK` register, or all other cases.

In the first case, the interrupt controller waits until the `INT_ACK` register has been written to. Then it looks at all the registered interrupts. The interrupt with the highest priority that has been asserted is recorded, becoming the next interrupt serviced by the ISR code. This process occurs regardless of whether a higher priority interrupt arrives between selection of an interrupt to service and when the ISR code is run.

In the second case, the next interrupt is whichever interrupt was registered. In the case of two registered interrupts appearing in the same clock cycle, the one with the highest priority is taken.

8.3.3 Interrupt Priorities

When the `INT_ACK` register is written to by the ISR code, and assuming the ISR code has support for nested interrupts, a higher priority interrupt can fire. The ISR code is rerun to handle the higher-priority interrupt. When the higher-priority ISR has completed, the previous interrupt is completed.

The ISR code should write to the `INT_ACK` register early on so that higher-priority interrupts do not wait long to be serviced.

The software ISR provides the nesting of interrupts (with state stored to the software stack). The `INT_LOAD_INFO` and `INT_SAVE_INFO` registers enable the state of the interrupt controller to be saved and restored, which is required for nested interrupts.

The `INT_LOAD_INFO` register also is used without a nested interrupt, because the `INT_LOAD_INFO_CLR_REQ` bit within this register clears the current interrupt from the store of registered interrupts, enabling further interrupts to occur.

With two interrupts at the same priority, the next interrupt (of same priority) can fire, after the `INT_LOAD_INFO` register has been written to. The same applies to a lower priority interrupt that follows. If a higher priority interrupt follows, then it can be serviced as soon as `INT_ACK` is written to.

8.3.4 DSP Core Functionality During Interrupt

When the DSP core receives an interrupt, an interrupt service routine is performed before the DSP can return to the routine running prior to the interrupt.

Upon reception of an interrupt (that is, when the interrupt controller decides to interrupt the Kalimba core because of a registered interrupt), the DSP core performs the following:

- Memory-mapped register `MM_RINTLINK` is loaded with the contents of current `PC`
- `PC` is loaded with the address of the interrupt service routine (for example, `0x0008`)
- The flags register is saved (that is, the flags in bits [0:7] are copied into bits [8:15]), and `UM_FLAG` is cleared
- If enabled, switches to a faster interrupt clock rate

When the interrupt service routine completes, the DSP must return to the routine that was running prior to the interrupt. This is executed with an `rti` instruction, which carries out the following:

- Restores the `rFlags` register to the noninterrupt state
- Loads `PC` with the contents of the `MM_RINTLINK` register

NOTE Saving and restoring of further registers is up to the programmer to design in software.

Further interrupts cannot occur while the interrupt service routine is executing, unless nested interrupts are enabled.

8.4 Generation of MCU Interrupt

Writing to the `DSP2MCU_EVENT_DATA` register sends an interrupt to the MCU, which can see the value of this register. Similarly, the MCU can generate an interrupt to the DSP. The event type is stored

in the `MCU2DSP_EVENT_DATA` register. These registers can pass messages between MCU and DSP and vice versa.

8.5 PIO Control from the Kalimba DSP

The Kalimba DSP can be used to control several PIO functions.

Control of the PIO from the Kalimba DSP is as follows:

- Kalimba DSP can read IC PIO lines and enable interrupts on them.
- Under the control of the MCU, the Kalimba DSP can write to PIO lines.
- Under the control of the MCU, the Kalimba DSP can change the direction of PIO lines.
- A PIO line change can generate a Kalimba DSP interrupt.

The MCU controls which PIO bits have write access permission for the Kalimba DSP. This information is set through VM functions, if present, or BCCMDs otherwise.

8.6 MCU Memory Windows in DM2

The two windows in DM2 enable the Kalimba DSP to access MCU memory. The primary use of this memory window is for passing message and control information.

The MCU controls access to this window. A start address and a size for each window can be set. Whether the DSP has read access or read and write access also can be set. The MCU firmware sets this information.

8.7 Nonvolatile Memory Windows in DM2

The nonvolatile memory windows in DM2 permit the DSP access to the flash or ROM (up to 32 Mb). This could be to access further coefficients or download a new program, for example.

The window size is 4 K words for the small windows and 1 M words for the large windows. The `FLASH_WINDOW[123]_START_ADDR` registers select which 4 K or 1 M block is visible to the DSP for each of the three windows.

8.8 PM Window in DM1

The PM window within the DM1 memory bank permits the DSP to change its own program or to use program memory as data memory, either in 16-bit words or 24-bit words. For safety, it is possible to disable the window.

The MCU must enable the DSP to have access to the PM mapped into DM1. The DSP must then enable access as well.

8.9 PM Nonvolatile Memory Window with Direct-mapped Cache

This window permits the DSP to access program memory directly from flash or ROM. DSP memory-mapped registers control the start address and size of this flash or ROM space.

The 32-bit DSP PM bus is mapped into the 16-bit flash or ROM space with alternate MS and LS words stored in flash or ROM.

The MCU must enable the DSP to have access to the PM nonvolatile memory window. The DSP also must enable access. When this interface is enabled, a 1024/512-word direct cache is mapped into the upper 512/256 words of physical program RAM, enabling zero-overhead accesses for cached instructions.

8.10 Clock Rate Divider Control

The Kalimba DSP can control its own clock frequency to run at 120 MHz, 80 MHz, and further subdivisions of 80 MHz.

8.11 Debugging

The MCU registers and the SPI can read and write any of the DSP core's internal registers, provide control (Single Step, Run, Stop, Breakpoints, and so forth), and read and write any individual location in PM, DM1, or DM2, as seen by the DSP.

9 Document References

Table 9-1 Document References

Document	Reference
<i>KAS Kalimba Architecture 5 DSP User Guide</i>	CS-00344866-UG
<i>Information technology -- Coding of moving pictures and associated audio for digital storage media at up to about 1,5 Mbit/s -- Part 3: Audio</i>	ISO/IEC 11172-3
<i>Information technology -- Generic coding of moving pictures and associated audio information -- Part 3: Audio</i>	ISO/IEC 13818-3
<i>Information technology -- Generic coding of moving pictures and associated audio information -- Part 7: Advanced Audio Coding (AAC)</i>	ISO/IEC 13818-7

10 Terms and Definitions

Term	Definition
A2DP	Advanced Audio Distribution Profile
AAC	Advanced Audio Coding
ADC	Analog to Digital Converter
AG	Address Generator
ALU	Arithmetic Logic Unit
BlueCore	Group term for the QTIL range of Bluetooth wireless technology ICs
Bluetooth	Set of technologies providing audio and data transfer over short-range radio connections
codec	Coder decoder
DAC	Digital to Analogue Converter
DSP	Digital Signal Processor (or Processing)
FFT	Fast Fourier Transform
FP	Frame Pointer
I/O	Input/Output
IC	Integrated Circuit
ISR	Interrupt Service Routine
Kalimba	An open platform DSP co-processor, enabling support of enhanced audio applications, such as echo and noise suppression and file compression / decompression
LS	Least Significant
LSB	Least Significant Bit (or Byte)
Mb	Megabit
MCU	MicroController Unit
MMU	Memory Management Unit
MP3	MPEG-1 audio layer 3
MS	Most Significant
MSB	Most Significant Bit (or Byte)
NOP	No OPeration
NVM	Non-Volatile Memory
PC	Program Counter
PIO	Programmable Input/Output, also known as general purpose I/O
RAM	Random Access Memory

Term	Definition
ROM	Read Only Memory
SBC	Sub-Band Coding
SP	Stack Pointer
SPI	Serial Peripheral Interface
VM	Virtual Machine

A Number Representation

This Appendix outlines the number representation used by the Kalimba DSP.

A.1 Binary Integer Representation

Kalimba uses 2's complement to represent signed integers.

Figure A-1 shows an example of this format, with only 8 bits shown, for clarity.




Figure A-1 Binary Integer Representation

A.2 Binary Fractional Representation

Kalimba implements a 24-bit fractional number type. The signed format can represent numbers from -1.0 to $+0.99999988$ (that is, $0x800000$ to $0x7fffff$). The unsigned format can represent numbers from 0 to 1.99999988 (that is, $0x000000$ to $0xffffffff$).

Figure A-2 shows a signed fractional number type. The Kalimba fractional type is 24 bits wide. The example shows only 8 bits, for clarity.


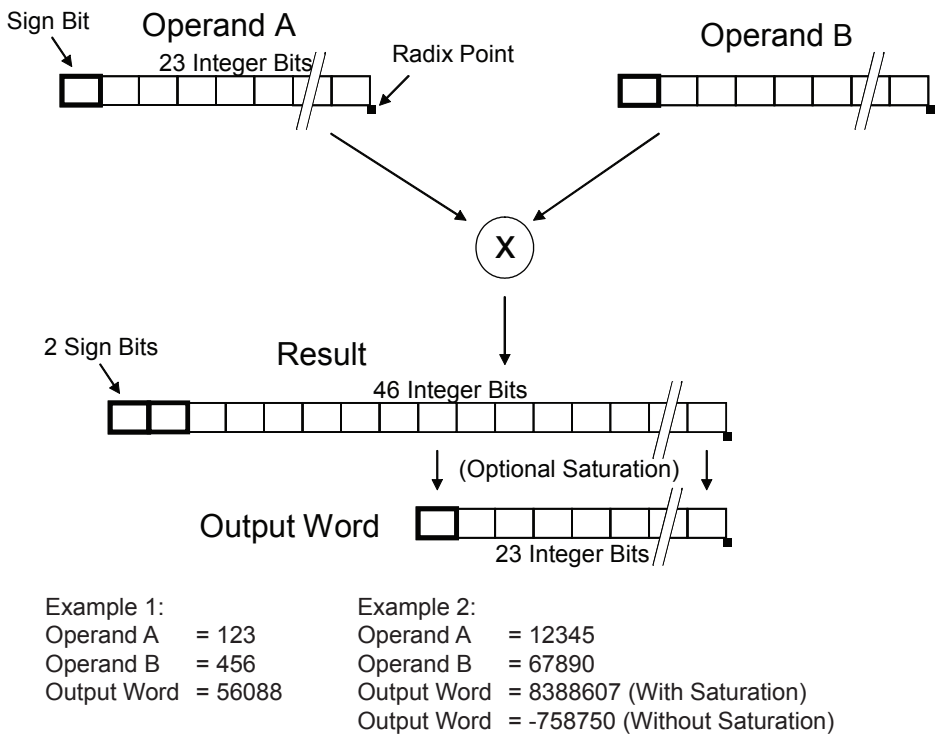


Figure A-2 Binary Fractional Representation

A.3 Integer Multiplication (Signed)



G-TW-0006701.2.2

Figure A-3 Integer Multiplication

A.4 Fractional Multiplication (Signed)

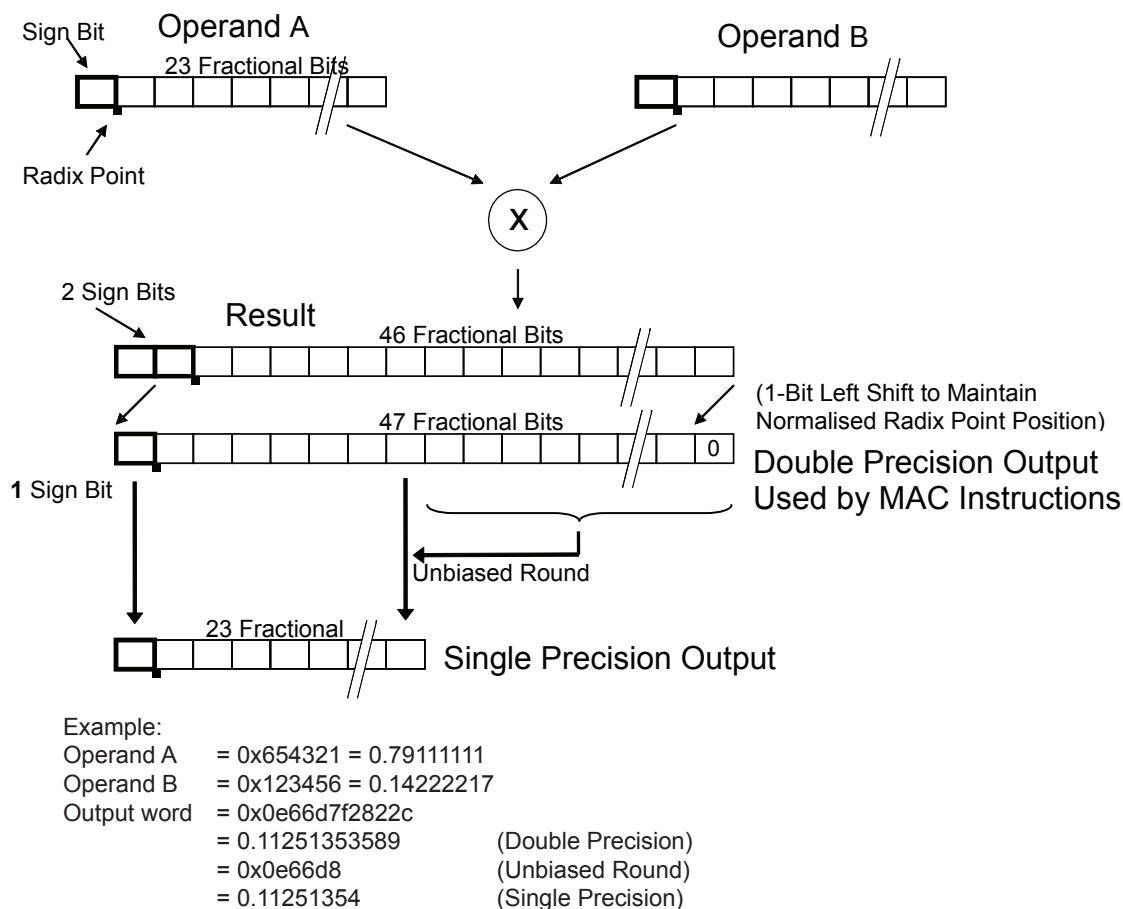


Figure A-4 Fractional Multiplication (Signed)

G-TW-0006702.1.2

B DSP Memory Mapped Registers

B.1 DSP Memory Mapped I/O

As described in Section 6.1.3, the DSP memory mapped I/O forms a reserved part of the DM2 memory map. See the appropriate `_io_map.h` header file in the ADK for the specific QTIL device for a list of the memory-mapped registers.

C Software Examples

This section contains samples of example code for the Kalimba DSP.

C.1 Double-Precision Addition

```
// Double-precision addition: Z = X + Y
//
// Where:
//   X = {r0,r1};           that is, r0 is MSW and r1 is LSW
//   Y = {r2,r3};
//   Z = {r5,r4};
// Computation time: 2 cycles

r4 = r1 + r3;               // add LSWs
r5 = r0 + r2 + Carry;      // add MSWs
```

C.2 Fractional Double-Precision Multiply

```
// Fractional double-precision multiply: Z = X * Y
//
// Where:
//   X = {r0,r1};           that is, r0 is MSW and r1 is LSW
//   Y = {r2,r3};
//   Z = {r4,r5,r6,r7};     that is, Z is 96-bit
// Computation time: 10 cycles

rMAC = r1 * r3 (UU);        // Compute LSW
r7 = rMAC0;                 // save Z0
rMAC = rMAC LSHIFT -24 (56bit); // shift right rMAC by 24 bits
rMAC = rMAC + r0 * r3 (SU);  // compute inner products
rMAC = rMAC + r2 * r1 (SU);
r6 = rMAC0;                 // save Z1
rMAC = rMAC ASHIFT -24 (56bit); // shift right rMAC by 24 bits
rMAC = rMAC + r0 * r2 (SS);  // compute MSWs
r5 = rMAC0;                 // save Z2
r4 = rMAC1;                 // save Z3
```

C.3 Integer Double-Precision Multiply

```
// Integer double-precision multiply: Z = X * Y
//
// Where:
//   X = {r0,r1};           that is, r0 is MSW and r1 is LSW
//   Y = {r2,r3};
//   Z = {r4,r5,r6,r7};     that is, Z is 96-bit
// Computation time: 10 cycles

rMAC = r1 * r3 (UU);           // Compute LSW
r7 = rMAC LSHIFT 23;          // save Z0
rMAC = rMAC LSHIFT -24 (56bit); // shift right 24 bits
rMAC = rMAC + r0 * r3 (SU);    // compute inner products
rMAC = rMAC + r2 * r1 (SU);
r6 = rMAC LSHIFT 23;          // save Z1
rMAC = rMAC ASHIFT -24 (56bit); // shift right 24 bits
rMAC = rMAC + r0 * r2 (SS);    // compute MSWs
r5 = rMAC LSHIFT 23;          // save Z2
r4 = rMAC LSHIFT -1;          // save Z3
```

C.4 FIR Filter

```
// FIR filter
//
// Input parameters:
//   I0 = points to oldest input value in delay line
//   L0 = filter length (N)
//   I4 = points to beginning of filter coefficient table
//   L4 = filter length (N)
//   r10 = filter length - 1 (N-1)
// Return values:
//   rMAC = sum of products output
// Computation time:
// N = 1: 5 cycles
// N = 2: 4 cycles
// N > 2: N + 6 cycles

fir_filter:
rMAC = 0,
  r1 = M[I0,1],
  r2 = M[I4,1];
do fir_loop;
  rMAC = rMAC + r1 * r2,
```

```

        r1 = M[I0,1],
        r2 = M[I4,1];
fir_loop:
rMAC = rMAC + r1 * r2;

```

C.5 Cascaded Bi-Quad IIR Filter

```

// Cascaded biquad IIR filter
//
// Equation of each section:
//  $y(n) = (b_0 \cdot x(n) + b_1 \cdot x(n-1) + b_2 \cdot x(n-2) - a_1 \cdot y(n-1) - a_2 \cdot y(n-2)) \ll \text{scalefactor}$ 
//
// Input Values:
//   r0 = input sample
//   I0 = points to oldest input value in delay line
//       (no biquads*2 + 2)
//   I1 = points to a list of scale factors for each biquad section
//   I4 = points to scaled coefficients b2,b1,b0,a2,a1,... etc
//   L0 = 2 * num_biquads + 2
//   L1 = num_biquads
//   L4 = 5 * num_biquads
//   M0 = -3
//   M1 = 1
//   r10 = num_biquads
// Return Values:
//   r0 = output sample
//   r10 - cleared
//   I0,I1,I4,L0,L1,L4,M0,M1 - unaffected
//   r1,r2,r3,r4 - affected
//
// Computation time:
// N = 1: 10 cycles
// N > 2: 8 * N + 5 cycles

biquad_filter:
do biquad_loop;
    r1 = M[I0,1], // get x(n-2)
    r2 = M[I4,1]; // get coef b2
    rMAC = r1 * r2,
    r3 = M[I0,1], // get x(n-1)
    r2 = M[I4,1]; // get coef b1
    rMAC = rMAC + r3 * r2,
    r4 = M[I1,1], // get scalefactor
    r2 = M[I4,1]; // get coef b0

```

```

    rMAC = rMAC + r0 * r2,
    r1 = M[I0,1],           // get y(n-2)
    r2 = M[I4,1];           // get coef a2
    rMAC = rMAC - r1 * r2,
    r1 = M[I0,M0],          // get y(n-1)
    r2 = M[I4,M1];          // get coef a1
    rMAC = rMAC - r1 * r2,
    M[I0,1] = r3;           // store new x(n-2)
    r0 = rMAC ASHIFT r4,
    M[I0,M1] = r0;          // store new x(n-1)
biquad_loop:
M[I0,1] = r1;              // store new y(n-2)
M[I0,1] = r0;              // store new y(n-1)

```

C.6 Radix-2 FFT

```

// An optimised FFT subroutine with a simple interface
//
// *****
// FILE
//      fast_fft.asm - a fast FFT routine
//
// CONTAINS
//      A fast optimised FFT routine with a simple interface
//
// fftnpts:    64    128    256    512    1024    2048
//
// No Clks:    TBA
//
// REGISTERS
//      On entry: L0,L1,L4,L5 should be initialised to 0.
//
//      On exit: All registers altered
//
// VARIABLES
//      Inputs:
//          fft_npts    - Number of points (a power of 2)
//          $Inputreal   - Input array real parts
//          $Inputimag   - Input array imag parts
//
//      Outputs:
//          $Refft       - Output array real parts
//          $Inputreal    - Output array imag parts
//

```

```

// *****

// Declare constants:

// Declare local variables:
.VAR/DM2 groups;
.VAR/DM2 node_space;
.include "twiddle_factors.h"

// Declare global variables:
.VAR/DM2 $fft_npts;

$fast_fft:

M1 = 1;
M2 = 2;
M3 = -1;

// -- Process the n-1 stages of butterflies --
r1 = 1;
M[groups] = r1;          // groups = 1
r0 = M[$fft_npts];
r0 = r0 ASHIFT -1;
M[node_space] = r0;      // node_space = Npts / 2

r0 = SIGNDDET r0;
r1 = 22;
r9 = r1 - r0;            // log2(Npts) - 1

stage_loop:

    r10 = M[node_space];
    M0 = r10;             // M0 = node_space

    r8 = M[groups];
    r2 = r8 LSHIFT 1;
    M[groups] = r2;       // groups = groups * 2;

    I2 = &$Inputreal;     // I2 -> x0 in 1st group of stage
    I4 = I2;
    I0 = &$Inputimag;     // I0 -> y0 in 1st group of stage
    I6 = I0;
    I1 = I2 + M0;          // I1 -> x1 in 1st group of stage
    I5 = I6 + M0;          // I3 -> y1 in 1st group of stage

    I3 = &twid_real;       // I3 -> C of W0
    I7 = &twid_imag;       // I7 -> (-S) of W0

```



```

group_loop:
    r10 = M0, r4 = M[I3,M1];
    r3 = M[I1,M1], r5 = M[I7,M1];
    r6 = r5;
    rMAC = r3*r4, r5 = M[I5,M1];
    rMAC = rMAC - r5 * r6, r1 = M[I2,M1], r2 = M[I6,M1];
    rMACB = r3*r6, r0 = r1+rMAC, r3=M[I1,M3];

    do bfly_loop;
        rMACB= rMACB+r5*r4, r0 = r1 - rMAC, r1 = M[I2,M1], M[I4,M1] =
r0;
        rMAC = r3*r4, r0 = r2 + rMACB, M[I1,M2] = r0, r5 =
M[I5,M3];
        rMAC = rMAC-r5*r6, r0 = r2 - rMACB, M[I0,M1] = r0, r2 =
M[I6,M1];
        rMACB= r3*r6, r0 = r1 + rMAC, r3 = M[I1,M3], M[I5,M2] =
r0;
    bfly_loop:

    I0 = I0 + M0, r0 = M[I1,M0];
    I5 = I5 + M0, r0 = M[I2,M0], r0 = M[I6,M0];
    I5 = I5 - M1, r0 = M[I2,M3], r0 = M[I6,M3];
    r8 = r8 - M1, r0 = M[I4,M0];
    if NZ jump group_loop;

    r10 = M[node_space];
    r10 = r10 ASHIFT -1; // node_space = node_space / 2;
    M[node_space] = r10;
    r9 = r9 - 1;
    if NZ jump stage_loop;

// -- Process the last stage of butterflies separately --
I0 = &twid_imag; // I0 -> (-S) of W0
I5 = &$Inputreal; // I2 -> x0
I1 = I5 + 1; // I1 -> x1
M2 = 2;
I3 = BITREVERSE(&$Refft); // Refft bitreversed
r0 = M[$fft_npts];
r0 = SIGNDET r0;
r0 = r0 + 1;
r1 = 1;
r1 = r1 LSHIFT r0;
M3 = r1; // Bitreversed modifier
M0 = 0;

I4 = &twid_real; // I4 -> C of W0
I6 = &$Inputimag; // I6 -> y0

```

```

I2 = I6 + 1;                                // I5 -> y1

r2 = M[I4,M1],
r5 = M[I2,M2];                             // r5=y1
r6 = r2,                                    // r6=C
r3 = M[I1,M2];                             // r3=x1

rMAC = r3 * r6,                             // rMAC=x1*C
r2 = M[I0,1];                              // r2=(-S)

r10 = M[$fft_npts];
r10 = r10 ASHIFT -1;                        // Npts / 2

do last_loop;
    rMAC = rMAC - r5 * r2,                  // rMAC=x1*C-y1*-S
    r0 = M[I5,M2];                         // r0=x0

    r1 = r0 + rMAC;                        // r1=x0'=x0+(x1*C-y1*-S)

    // enable Bit Reverse addressing on AG1
    rFlags = rFlags OR $BR_FLAG;

    r1 = r0 - rMAC,                        // r1=x1'=x0-(x1*C-y1*-S)
    M[I3,M3] = r1;                        // DM=x0'

    rMAC = r3 * r2,                        // rMAC=x1*(-S)
    M[I3,M3] = r1,                        // DM=x1'
    r4 = M[I6,M0];                        // r4=y0

    // disable Bit Reverse addressing on AG1
    rFlags = rFlags AND $NOT_BR_FLAG;

    rMAC = rMAC + r5 * r6,                 // rMAC=x1*(-S)+y1*C
    r2 = M[I4,M1],                         // r2=C;
    r3 = M[I1,M2];                         // r3=next x1

    r1 = r4 + rMAC,                        // r1=y0'=y0+(y1*C+x1*(-S))
    r5 = M[I2,M2];                         // r5=next y1

    r4 = r4 - rMAC,                        // r1=y1'=y0-(y1*C+x1*(-S))
    M[I6,M1] = r1;                        // DM=y0'

    r6 = r2,                              // r6=C
    r2 = M[I0,M1];                         // r2=(-S)

    rMAC = r3 * r6,                        // rMAC=x1*C
    M[I6,M1] = r4;                        // DM=y1'

```

```
last_loop:

I3 = BITREVERSE(&$Inputreal);

I5 = &$Inputimag;

// enable Bit Reverse addressing on AG1
rFlags = rFlags OR $BR_FLAG;

r2 = M[I5,1];
r10 = $N;
DO bit_rev_imag;
    r2 = M[I5,M1],
    M[I3,M3] = r2;
bit_rev_imag:

// disable Bit Reverse addressing on AG1
rFlags = rFlags AND $NOT_BR_FLAG;
rts;
```

D Bit-reversed Addressing Explained

Bit-reverse addressing is a processor mode designed to improve the performance of radix-2 FFT implementations. When enabled, the address that is fed out of the DSP core to the physical memory is *bit-reversed* just before it connects to the memory.

The mode can be enabled only on the addresses being formed by what is called *address generator 1*. This is the address generator concerned with index memory accesses using I0 to I3. The mode is enabled by using the BR_FLAG bit in the rFlags register (see Section 4.5).

When bit-reversed mode is enabled, the contents of index registers appear incorrect. However, these address values are reversed again before going to physical memory. For example, an 8-word buffer located in physical memory at address 0xff8400.

To access this buffer with bit-reversed mode enabled, access the first element of the buffer using the following code:

```
I0 = BITREVERSE(0xff8400);  
r0 = M[I0,0];
```

Here, I0 equals 0x8010ff. When bit-reversed again, I0 becomes 0xff8400.

Bit-reversing on Kalimba Architecture 5 consists of reversing bits[0:22], but leaving bit[23] in place. This is because bit[23] selects the memory bank, DM1 or DM2.

Bit-reversed addressing enables processing through a buffer in what is called *bit-reversed order*, as opposed to the simpler linear order. If the buffer is of size 8, when processed in linear order, the results are offsets of 0, 1, 2, 3, 4, 5, 6, and 7. However, when processed in bit-reversed order, the offsets are 0, 4, 2, 6, 1, 5, 3, and 7. This order is worked out by looking at the binary representation of the offsets.

Table D-1 Binary Representation of 8-Word Buffer

Decimal Equivalent	Binary	Bit-reversed Binary	Decimal (Bit-reversed)
0	000	000	0
1	001	100	4
2	010	010	2
3	011	110	6
4	100	001	1
5	101	101	5
6	110	011	3
7	111	111	7

To process the buffer in bit-reversed order means accessing the physical addresses.

- 0xff8400
- 0xff8404
- 0xff8402
- 0xff8406
- 0xff8401
- 0xff8405
- 0xff8403
- 0xff8407

With bit-reversed mode enabled, the DSP address register (for example, `I0`) must point to the following:

- 0x8010ff
- 0x9010ff
- 0xa010ff
- 0xb010ff
- 0xc010ff
- 0xd010ff
- 0xe010ff
- 0xf010ff

Bit reversed addressing is achieved by initialising `I0` to 0x8010ff and then setting a modify register (for example, `M0`) to 0x100000.

For the more generic case of a buffer of size N (where N is always a power of 2), the modify register must be set to a value of $2^{(23-\log_2(N))}$.